

# Secure Systems Engineering

## Program Binaries

Chester Rebeiro

Indian Institute of Technology Madras

# Executables and Processes

```
#include <stdio.h>

int main(){
    char str[] = "Hello World\n";
    printf("%s", str);
}
```

\$gcc hello.c

Executable  
ELF  
(a.out)

./a.out

Process

Stored in disk



Executed from RAM

# ELF Executables

```
#include <stdio.h>

int main(){
    char str[] = "Hello World\n";
    printf("%s", str);
}
```

\$gcc hello.c

Executable  
ELF  
(a.out)

./a.out

Process



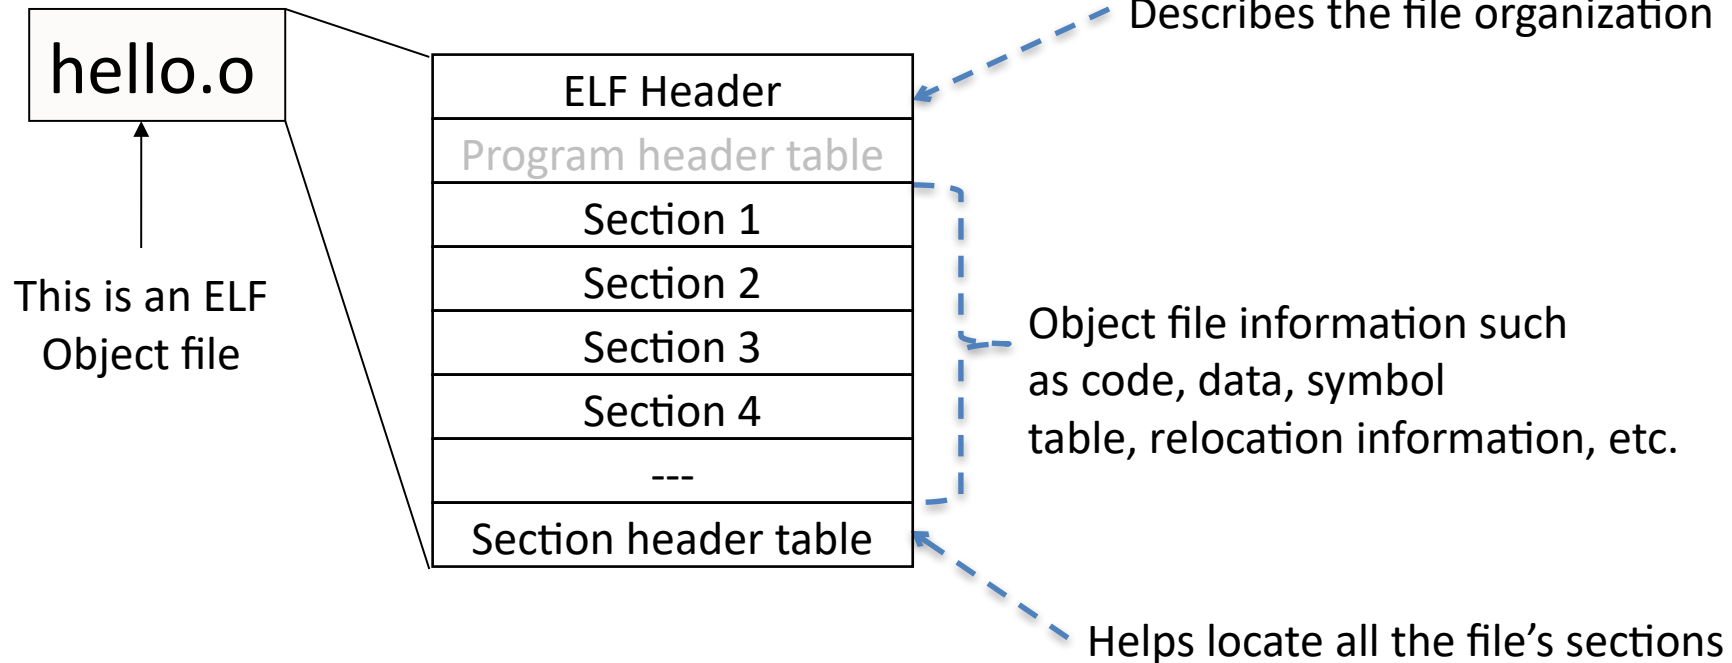
Stored in disk



Executed from RAM

# ELF Executables (linker view)

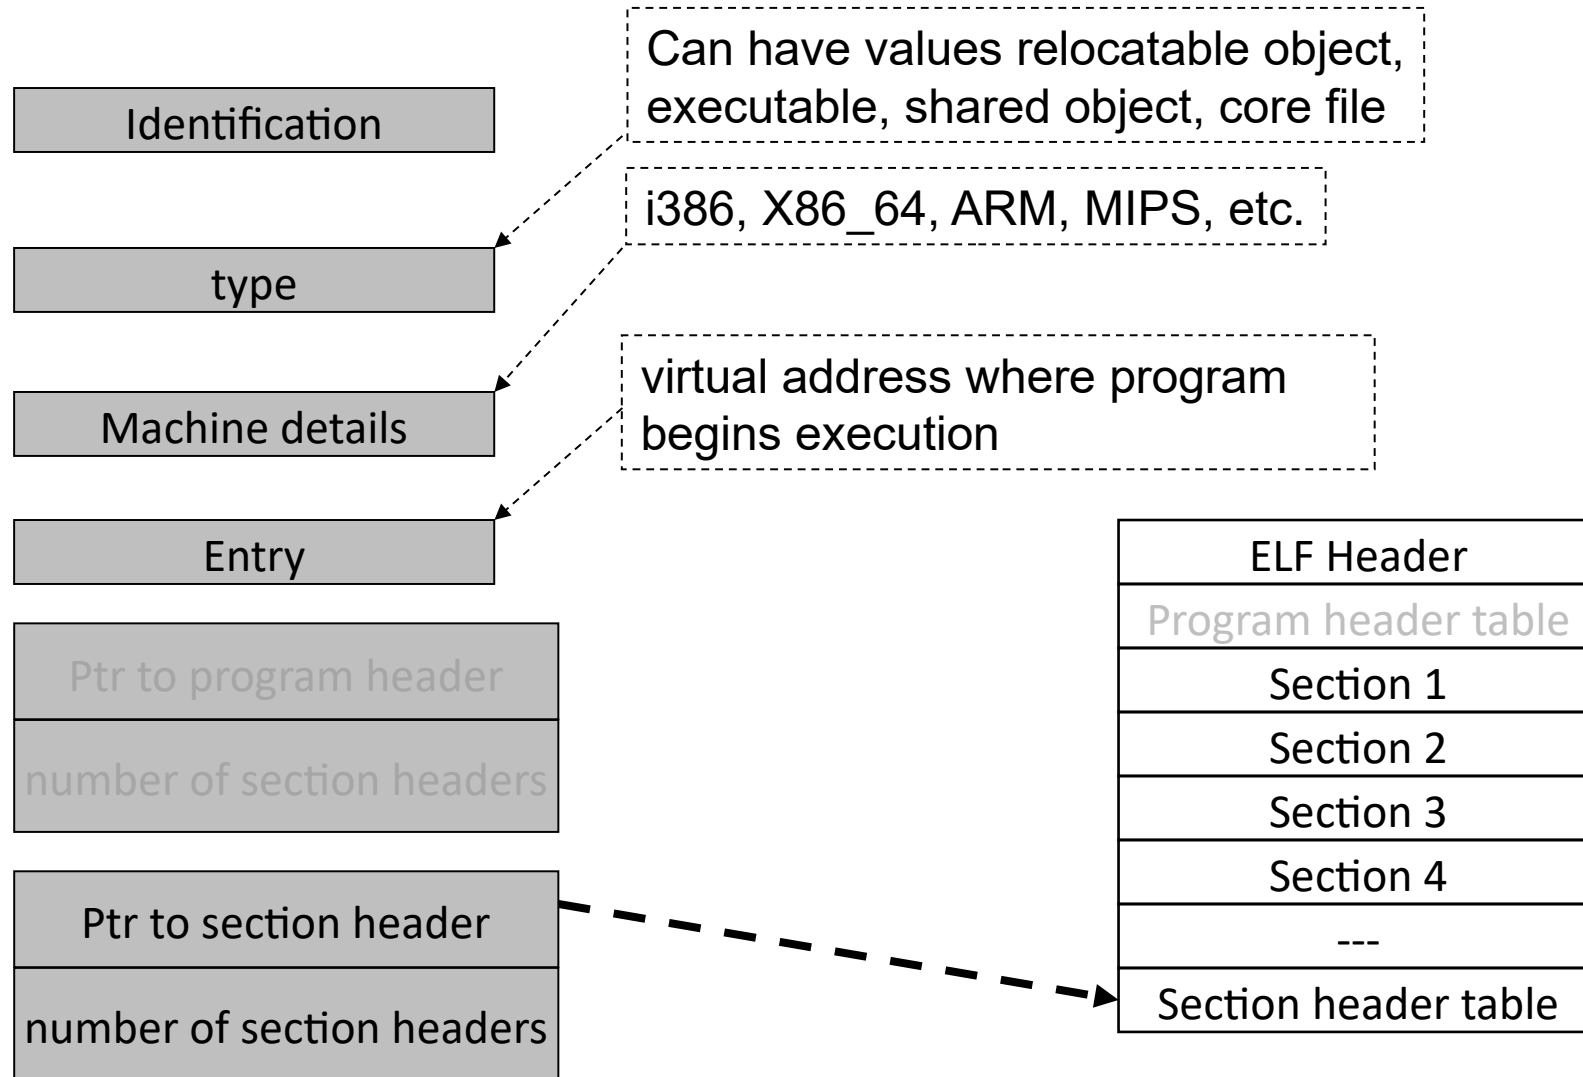
\$gcc hello.c -c



ref :[www.skyfree.org/linux/references/ELF\\_Format.pdf](http://www.skyfree.org/linux/references/ELF_Format.pdf)

ref :man elf

# ELF Header



# Hello World's ELF Header

```
#include <stdio.h>

int main(){
    char str[] = "Hello World\n";
    printf("%s", str);
}
```

```
$ gcc hello.c -c
$ readelf -h hello.o
```

```
chester@optiplex:~/tmp$ readelf -h hello.o
ELF Header:
  Magic:   7f 45 4c 46 02 01 01 00 00 00 00 00 00 00 00 00
  Class:                               ELF64
  Data:                                   2's complement, little endian
  Version:                             1 (current)
  OS/ABI:                               UNIX - System V
  ABI Version:                          0
  Type:                                 REL (Relocatable file)
  Machine:                              Advanced Micro Devices X86-64
  Version:                              0x1
  Entry point address:                   0x0
  Start of program headers:              0 (bytes into file)
  Start of section headers:             368 (bytes into file)
  Flags:                                 0x0
  Size of this header:                   64 (bytes)
  Size of program headers:               0 (bytes)
  Number of program headers:             0
  Size of section headers:               64 (bytes)
  Number of section headers:             13
  Section header string table index:    10
```

# Section Headers

Contains  
information about  
the various sections

\$ readelf -S hello.o

```
chester@optiplex:~/work/SSE/sse/src/elf$ readelf -S hello.o
There are 13 section headers, starting at offset 0x138:

Section Headers:
 [Nr] Name              Type            Addr           Off          Size      ES Flg Lk  Inf Al
 [ 0]                   NULL            00000000      000000      000000      00   0  0  0  0
 [ 1] .text                PROGBITS        00000000      000034      00003c      00  AX  0  0  1
 [ 2] .rel.text            REL             00000000      000408      000010      08   11  1  4
 [ 3] .data                PROGBITS        00000000      000070      000000      00  WA  0  0  1
 [ 4] .bss                 NOBITS          00000000      000070      000000      00  WA  0  0  1
 [ 5] .rodata              PROGBITS        00000000      000070      000003      00   A  0  0  1
 [ 6] .comment              PROGBITS        00000000      000073      00002c      01  MS  0  0  1
 [ 7] .note.GNU-stack       PROGBITS        00000000      00009f      000000      00   0  0  1
 [ 8] .eh_frame             PROGBITS        00000000      0000a0      000038      00   A  0  0  4
 [ 9] .rel.eh_frame         REL             00000000      000418      000008      08   11  8  4
[10] .shstrtab             STRTAB          00000000      0000d8      00005f      00   0  0  1
[11] .symtab               SYMTAB          00000000      000340      0000b0      10   12  9  4
[12] .strtab               STRTAB          00000000      0003f0      000015      00   0  0  1

Key to Flags:
 W (write), A (alloc), X (execute), M (merge), S (strings)
 T (info), L (link order), G (group), T (TLS), E (exclude), x (unknown)
 0 (extra OS processing required) o (OS specific), p (processor specific)
chester@optiplex:~/work/SSE/sse/src/elf$
```

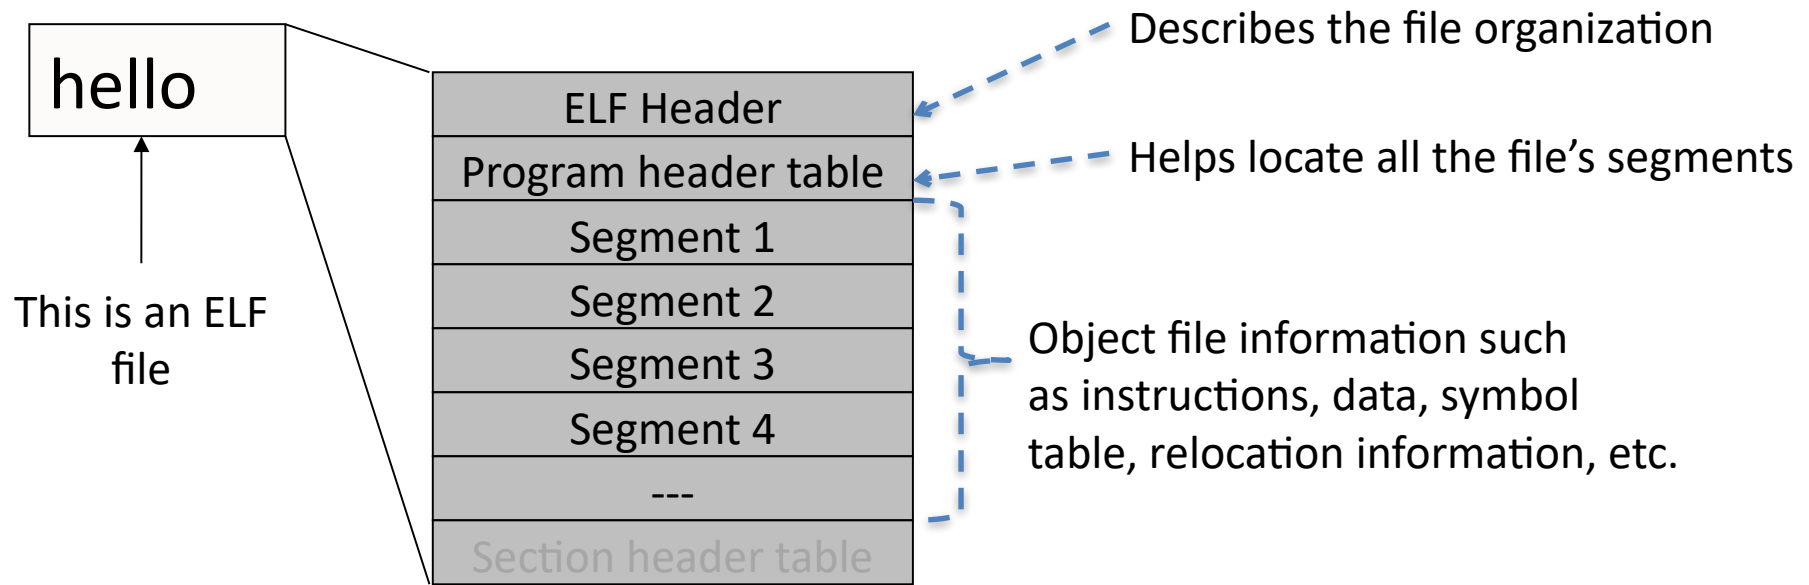
Type of the section  
PROGBITS : information defined by program  
SYMTAB : symbol table  
NULL : inactive section  
NOBITS : Section that occupies no bits  
RELA : Relocation table

Virtual address where the  
Section should be loaded  
(\* all 0s because this is a .o file)

Offset and size of the section

# ELF Executables (Executable view)

```
$gcc hello.c -o hello
```



ELF Executable View


ref :[www.skyfree.org/linux/references/ELF\\_Format.pdf](http://www.skyfree.org/linux/references/ELF_Format.pdf)

ref :man elf



# Program Header (executable view)

- Contains information about each segment
- One program header for each segment



# Program Header Contents

|                    |  |
|--------------------|--|
| type               | → type of segment (loadable segment / shared lib /etc)             |
| offset             | → Offset of segment in ELF file                                    |
| vaddr offset       | → Virtual address where the segment is to be loaded                |
| paddr offset       | → physical address where the segment is to be loaded.<br>(ignored) |
| Size in file image |  |
| Size in memory     |  |
| flags              | → Read / write / executable  |

# Program headers for Hello World

\$ readelf -l hello

```
chester@optiplex:~/work/SSE/sse/src/elf$ readelf -l hello

Elf file type is EXEC (Executable file)
Entry point 0x8048320
There are 9 program headers, starting at offset 52

Program Headers:
  Type           Offset   VirtAddr   PhysAddr   FileSiz MemSiz  Flg Align
  PHDR           0x000034 0x08048034 0x08048034 0x00120 0x00120 R E 0x4
  INTERP        0x000154 0x08048154 0x08048154 0x00013 0x00013 R   0x1
      [Requesting program interpreter: /lib/ld-linux.so.2]
  LOAD          0x000000 0x08048000 0x08048000 0x005d0 0x005d0 R E 0x1000
  LOAD          0x000f08 0x08049f08 0x08049f08 0x00118 0x0011c RW 0x1000
  DYNAMIC        0x000f14 0x08049f14 0x08049f14 0x000e8 0x000e8 RW 0x4
  NOTE          0x000168 0x08048168 0x08048168 0x00044 0x00044 R   0x4
  GNU_EH_FRAME   0x0004f4 0x080484f4 0x080484f4 0x0002c 0x0002c R   0x4
  GNU_STACK      0x000000 0x00000000 0x00000000 0x00000 0x00000 RWE 0x10
  GNU_RELRO      0x000f08 0x08049f08 0x08049f08 0x000f8 0x000f8 R   0x1

Section to Segment mapping:
Segment Sections...
 00
 01      .interp
 02      .interp .note.ABI-tag .note.gnu.build-id .gnu.hash .dynsym .dynstr
 03      .rodata .eh_frame_hdr .eh_frame
 04      .init_array .fini_array .jcr .dynamic .got .got.plt .data .bss
 05      .dynamic
 06      .note.ABI-tag .note.gnu.build-id
 07      .eh_frame_hdr
 08      .init_array .fini_array .jcr .dynamic .got
```

Mapping between  
segments and sections

# Contents of the Executable

\$ objdump --disassemble-all hello > hello.lst

```
#include <stdio.h>
```

```
int main(){  
    char str[] = "Hello World\n";  
    printf("%s", str);  
}
```

```
0804841d <main>:  
804841d: 55          push    %ebp  
804841e: 89 e5       mov     %esp,%ebp  
8048420: 83 e4 f0    and     $0xfffffffff0,%esp  
8048423: 83 ec 20    sub     $0x20,%esp  
8048426: c7 44 24 13 48 65 6c movl    $0x6c6c6548,0x13(%esp)  
804842d: 6c           
804842e: c7 44 24 17 6f 20 57 movl    $0x6f57206f,0x17(%esp)  
8048435: 6f           
8048436: c7 44 24 1b 72 6c 64 movl    $0xa646c72,0x1b(%esp)  
804843d: 0a           
804843e: c6 44 24 1f 00 movb    $0x0,0x1f(%esp)  
8048443: 8d 44 24 13 lea     0x13(%esp),%eax  
8048447: 89 44 24 04 mov     %eax,0x4(%esp)  
804844b: c7 04 24 f0 84 04 08 movl    $0x80484f0,(%esp)  
8048452: e8 99 fe ff ff call    80482f0 <printf@plt>  
8048457: c9         leave     
8048458: c3         ret       
8048459: 66 90      xchg    %ax,%ax  
804845b: 66 90      xchg    %ax,%ax  
804845d: 66 90      xchg    %ax,%ax  
804845f: 90         nop
```

# ELF Executables

```
#include <stdio.h>

int main(){
    char str[] = "Hello World\n";
    printf("%s", str);
}
```

\$gcc hello.c

Executable  
ELF  
(a.out)

Process

./a.out




Stored in disk



Executed from RAM

# Creating a Process by Cloning (using fork system call)



```
int p;  
  
p = fork();  
if (p > 0){  
    printf("Parent : child PID = %d", p);  
    p = wait();  
    printf("Parent : child %d exited\n", p);  
} else{  
    printf("In child process");  
    execlp("hello", "", NULL);  
    exit(0);  
}
```

# Process Virtual Memory Map

```
#include <stdio.h>
#include <stdlib.h>

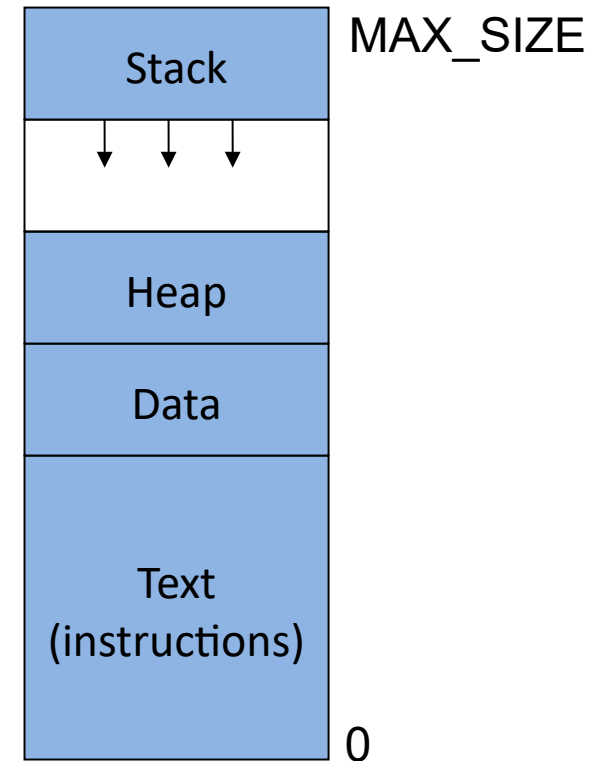
int calls;

void fact(int a, int *b){
    calls++;
    if (a==1) return;
    *b = *b * a;
    fact(a - 1, b);
}

int main(){
    int n, *m;

    scanf("%d", &n);
    m = malloc(sizeof(int));
    *m = 1;
    fact(n, m);
    printf("Factorial(%d) is %d\n", n, *m);
    free(m);
}
```

Program



Virtual Memory Map

# Process Virtual Memory Map

```
chester@optiplex:~$ ps -ae | grep hello
6757 pts/25 00:00:00 hello
chester@optiplex:~$ sudo cat /proc/6757/maps
08048000-08049000 r-xp 00000000 08:07 2491006 /home/chester/work/SSE/sse/src/elf/hello
08049000-0804a000 r-xp 00000000 08:07 2491006 /home/chester/work/SSE/sse/src/elf/hello
0804a000-0804b000 rwxp 00001000 08:07 2491006 /home/chester/work/SSE/sse/src/elf/hello
f759f000-f75a0000 rwxp 00000000 00:00 0
f75a0000-f774b000 r-xp 00000000 08:06 280150 /lib/i386-linux-gnu/libc-2.19.so
f774b000-f774d000 r-xp 001aa000 08:06 280150 /lib/i386-linux-gnu/libc-2.19.so
f774d000-f774e000 rwxp 001ac000 08:06 280150 /lib/i386-linux-gnu/libc-2.19.so
f774e000-f7751000 rwxp 00000000 00:00 0
f7773000-f7777000 rwxp 00000000 00:00 0
f7777000-f7778000 r-xp 00000000 00:00 0 [vdso]
f7778000-f7798000 r-xp 00000000 08:06 280158 /lib/i386-linux-gnu/ld-2.19.so
f7798000-f7799000 r-xp 0001f000 08:06 280158 /lib/i386-linux-gnu/ld-2.19.so
f7799000-f779a000 rwxp 00020000 08:06 280158 /lib/i386-linux-gnu/ld-2.19.so
ff885000-ff8a6000 rwxp 00000000 00:00 0 [stack]
chester@optiplex:~$
```

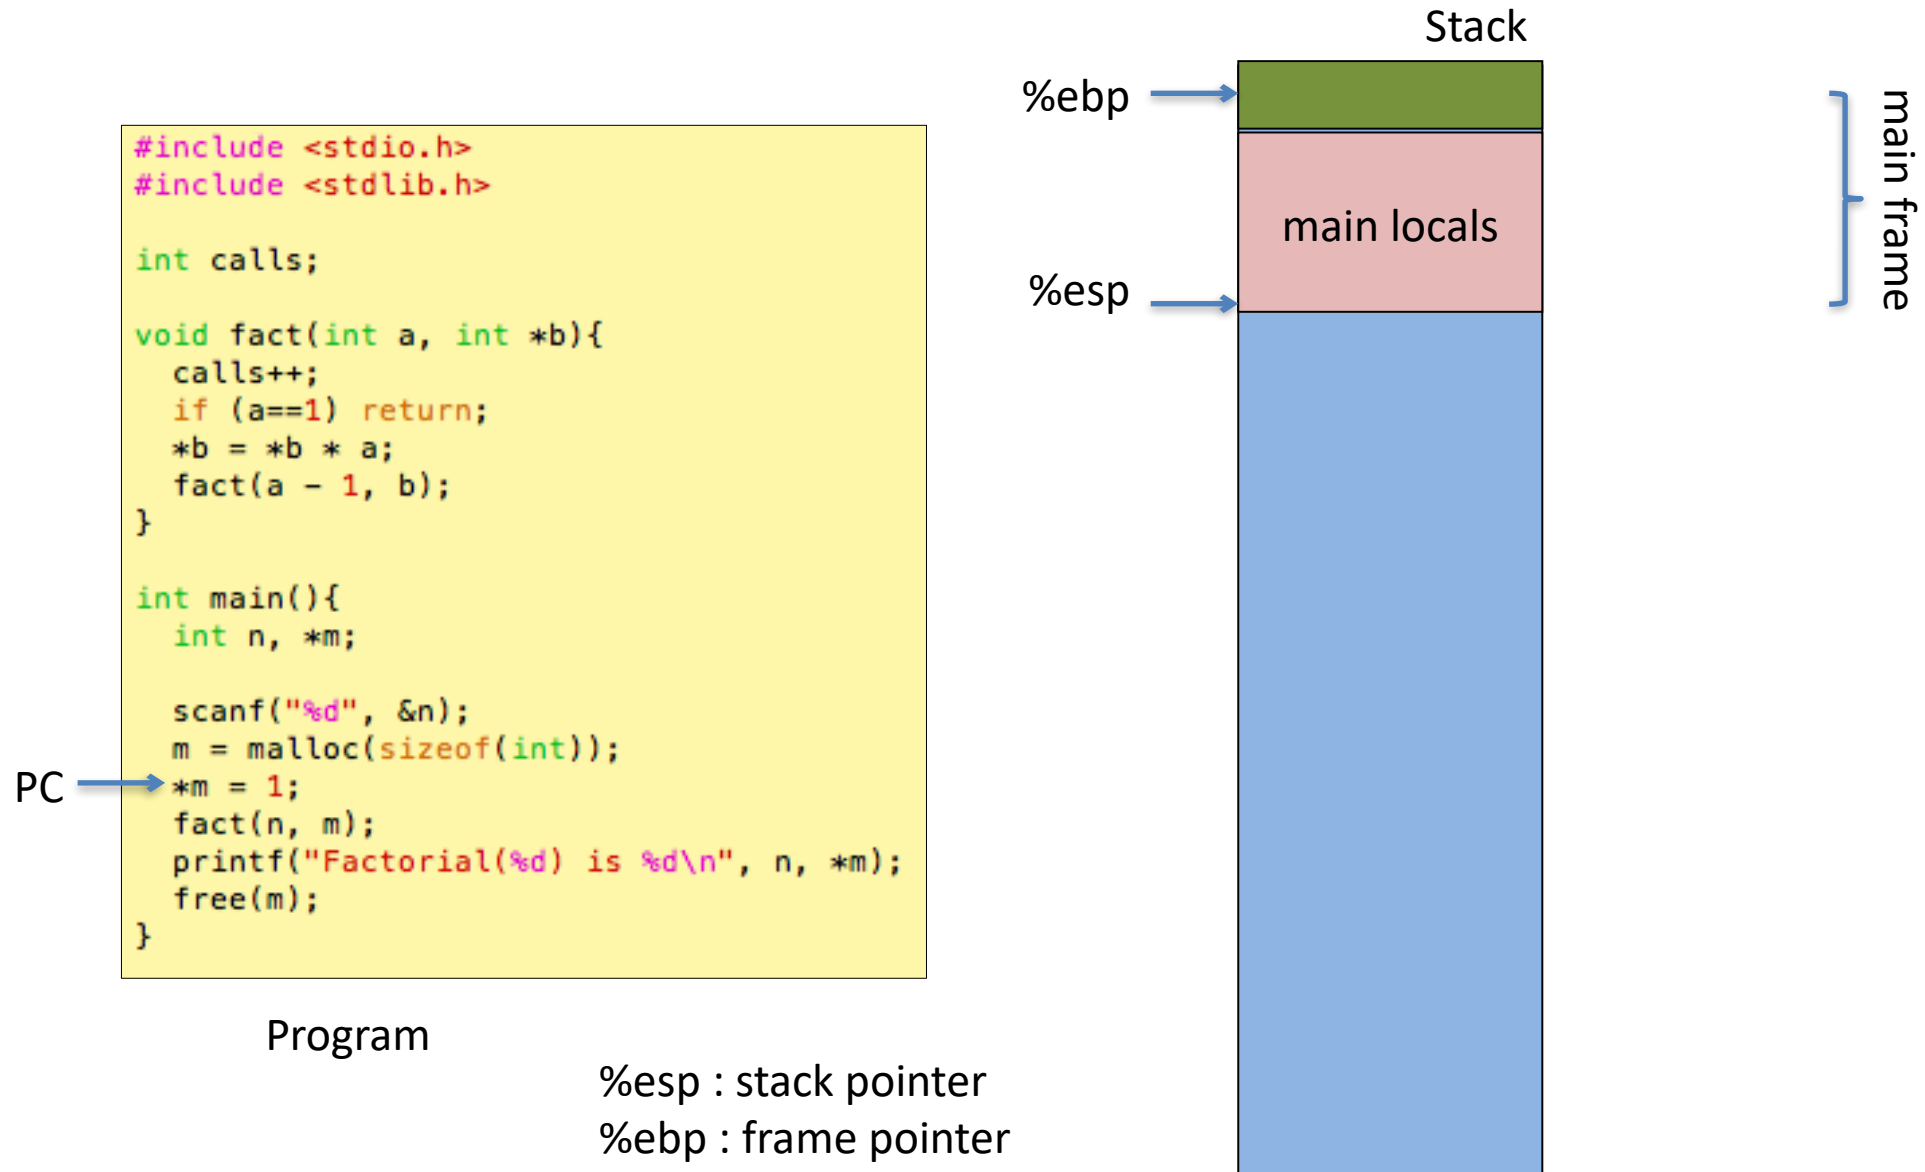
Virtual address  
memory range

flags

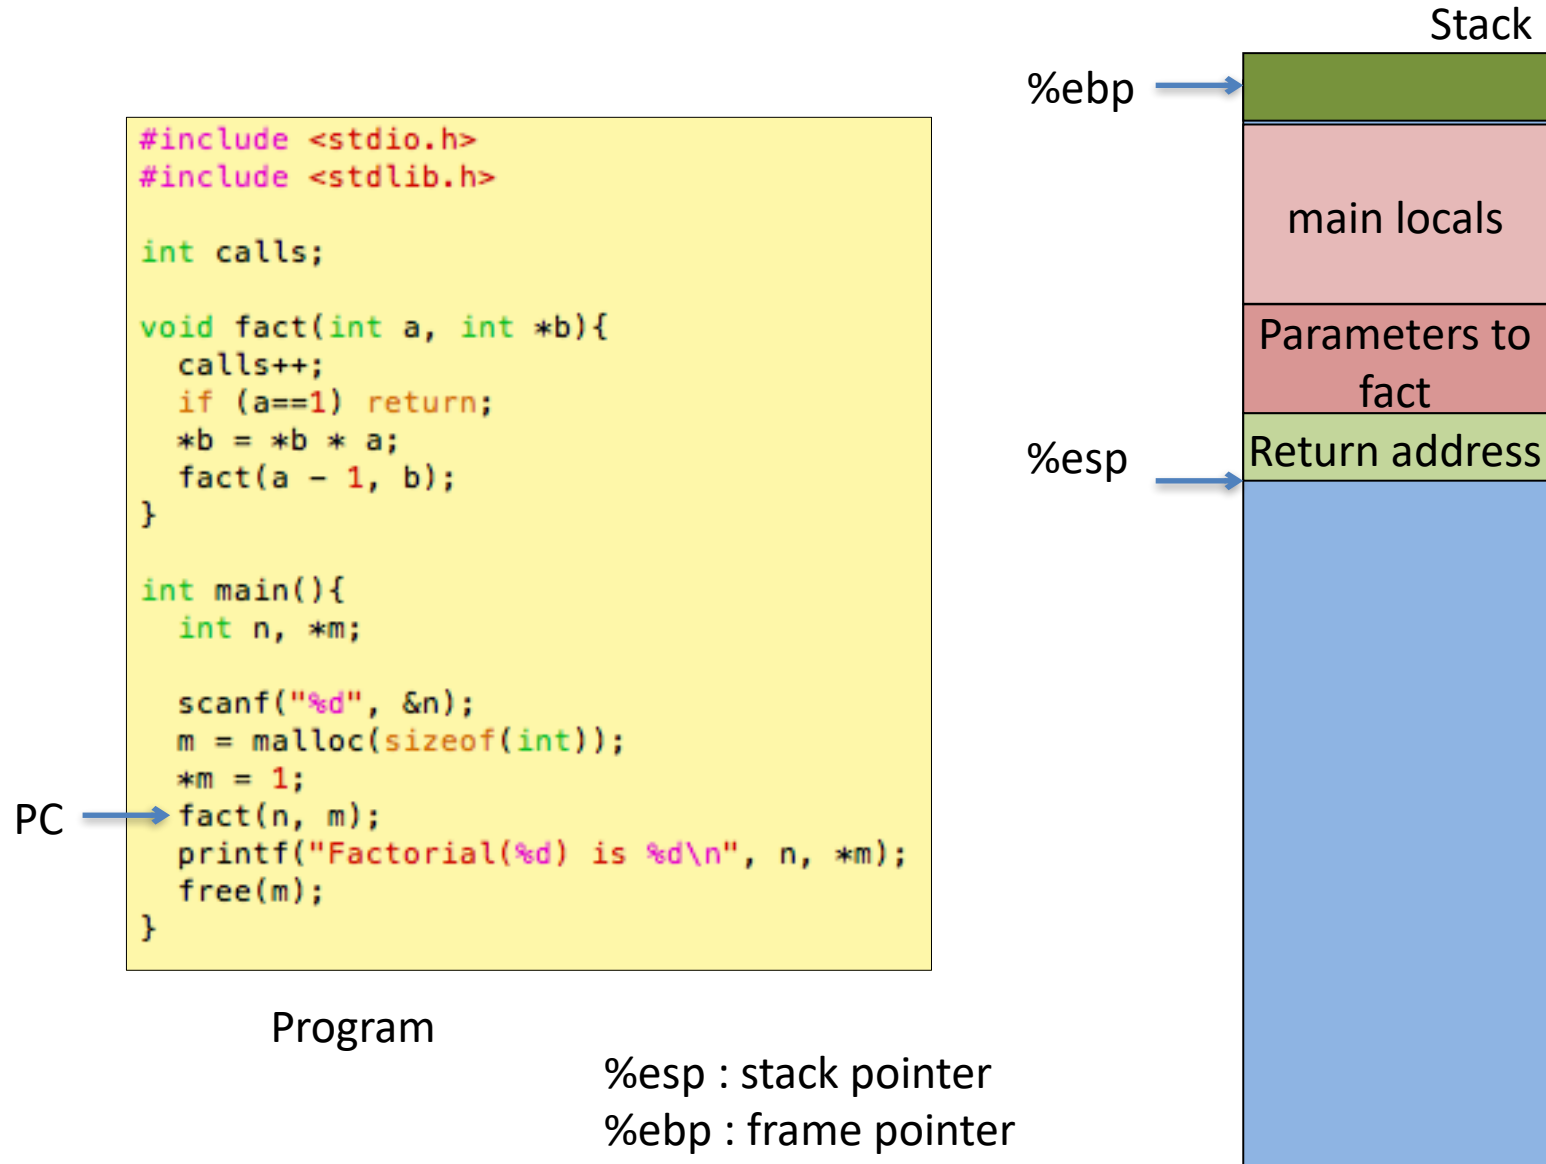
Device details  
(offset in file; device number; inode)



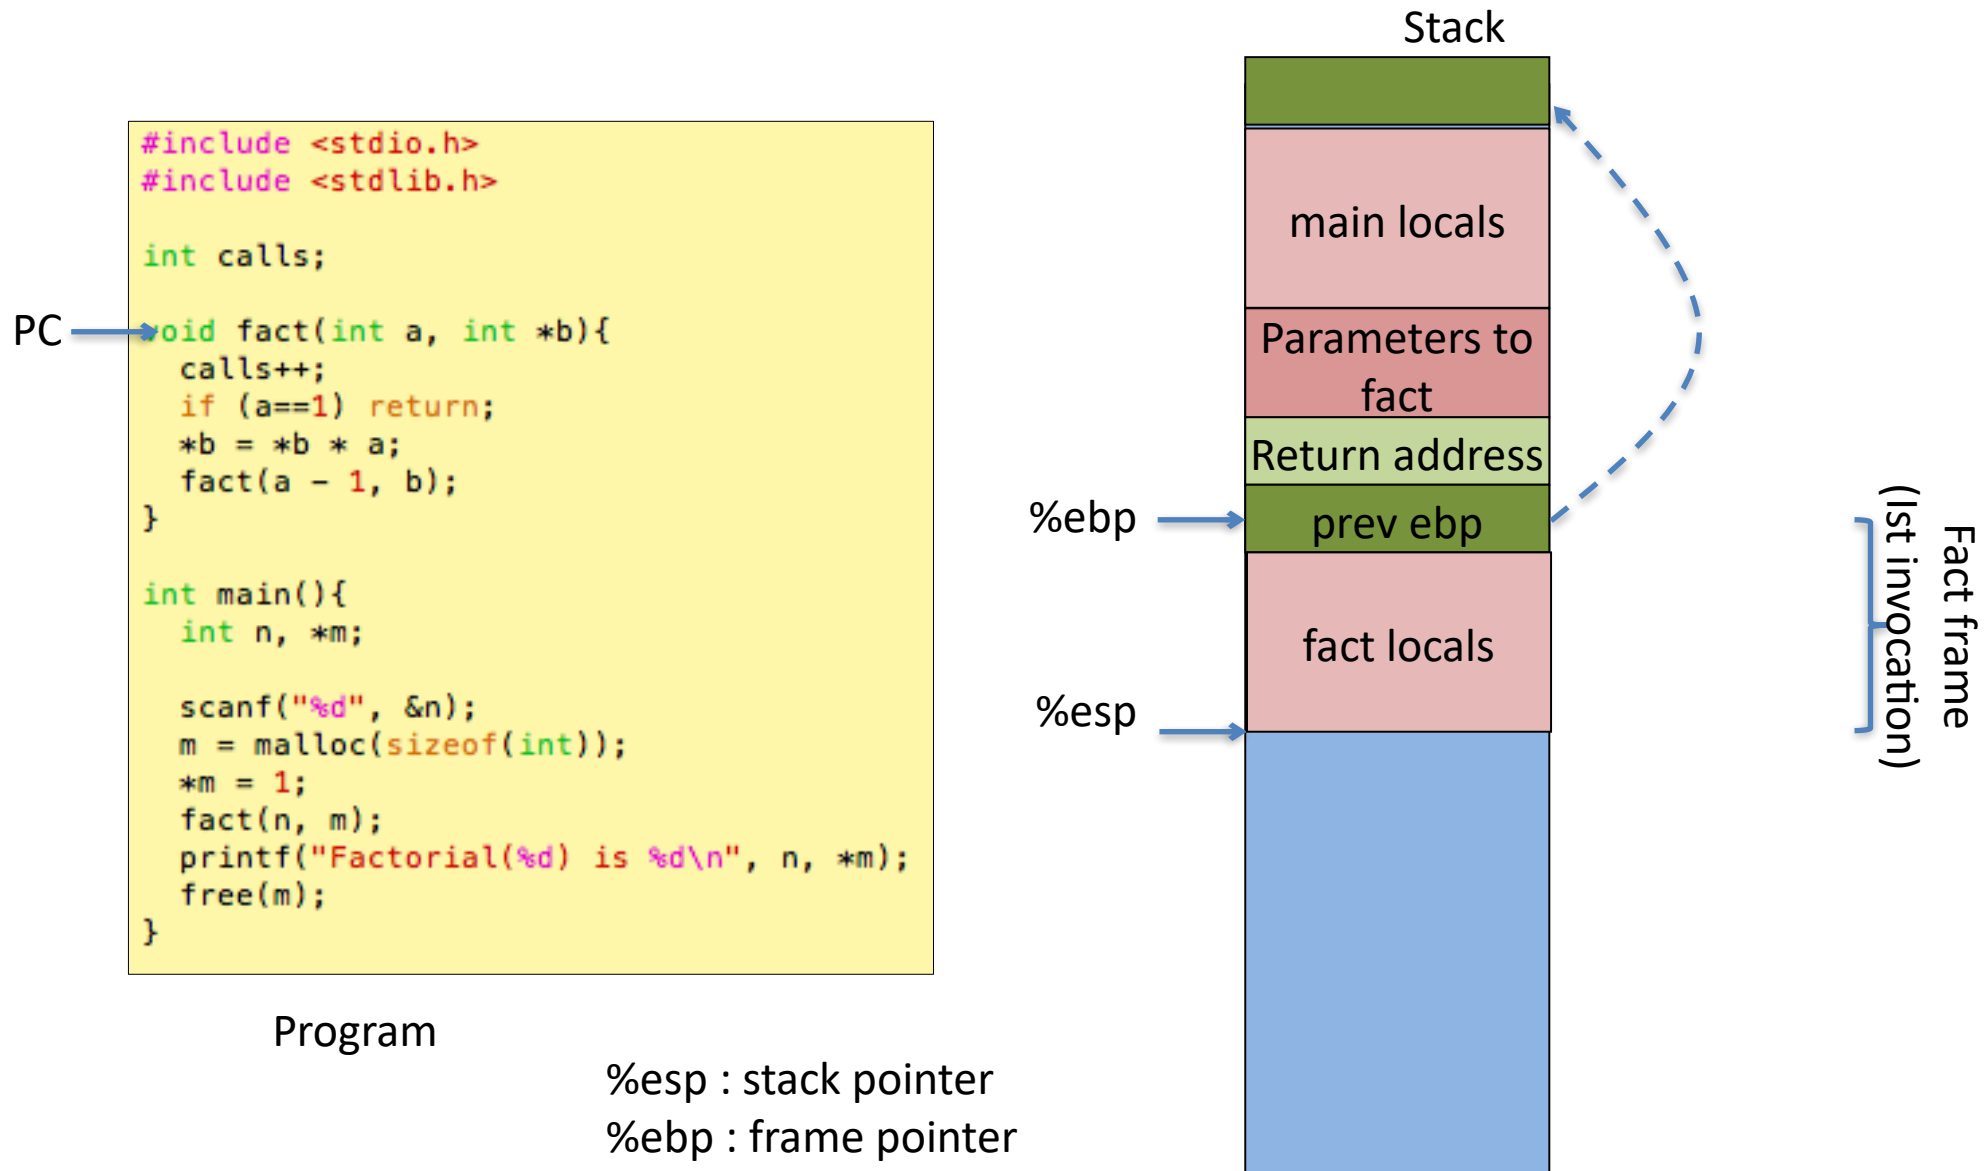
# Stack Frames



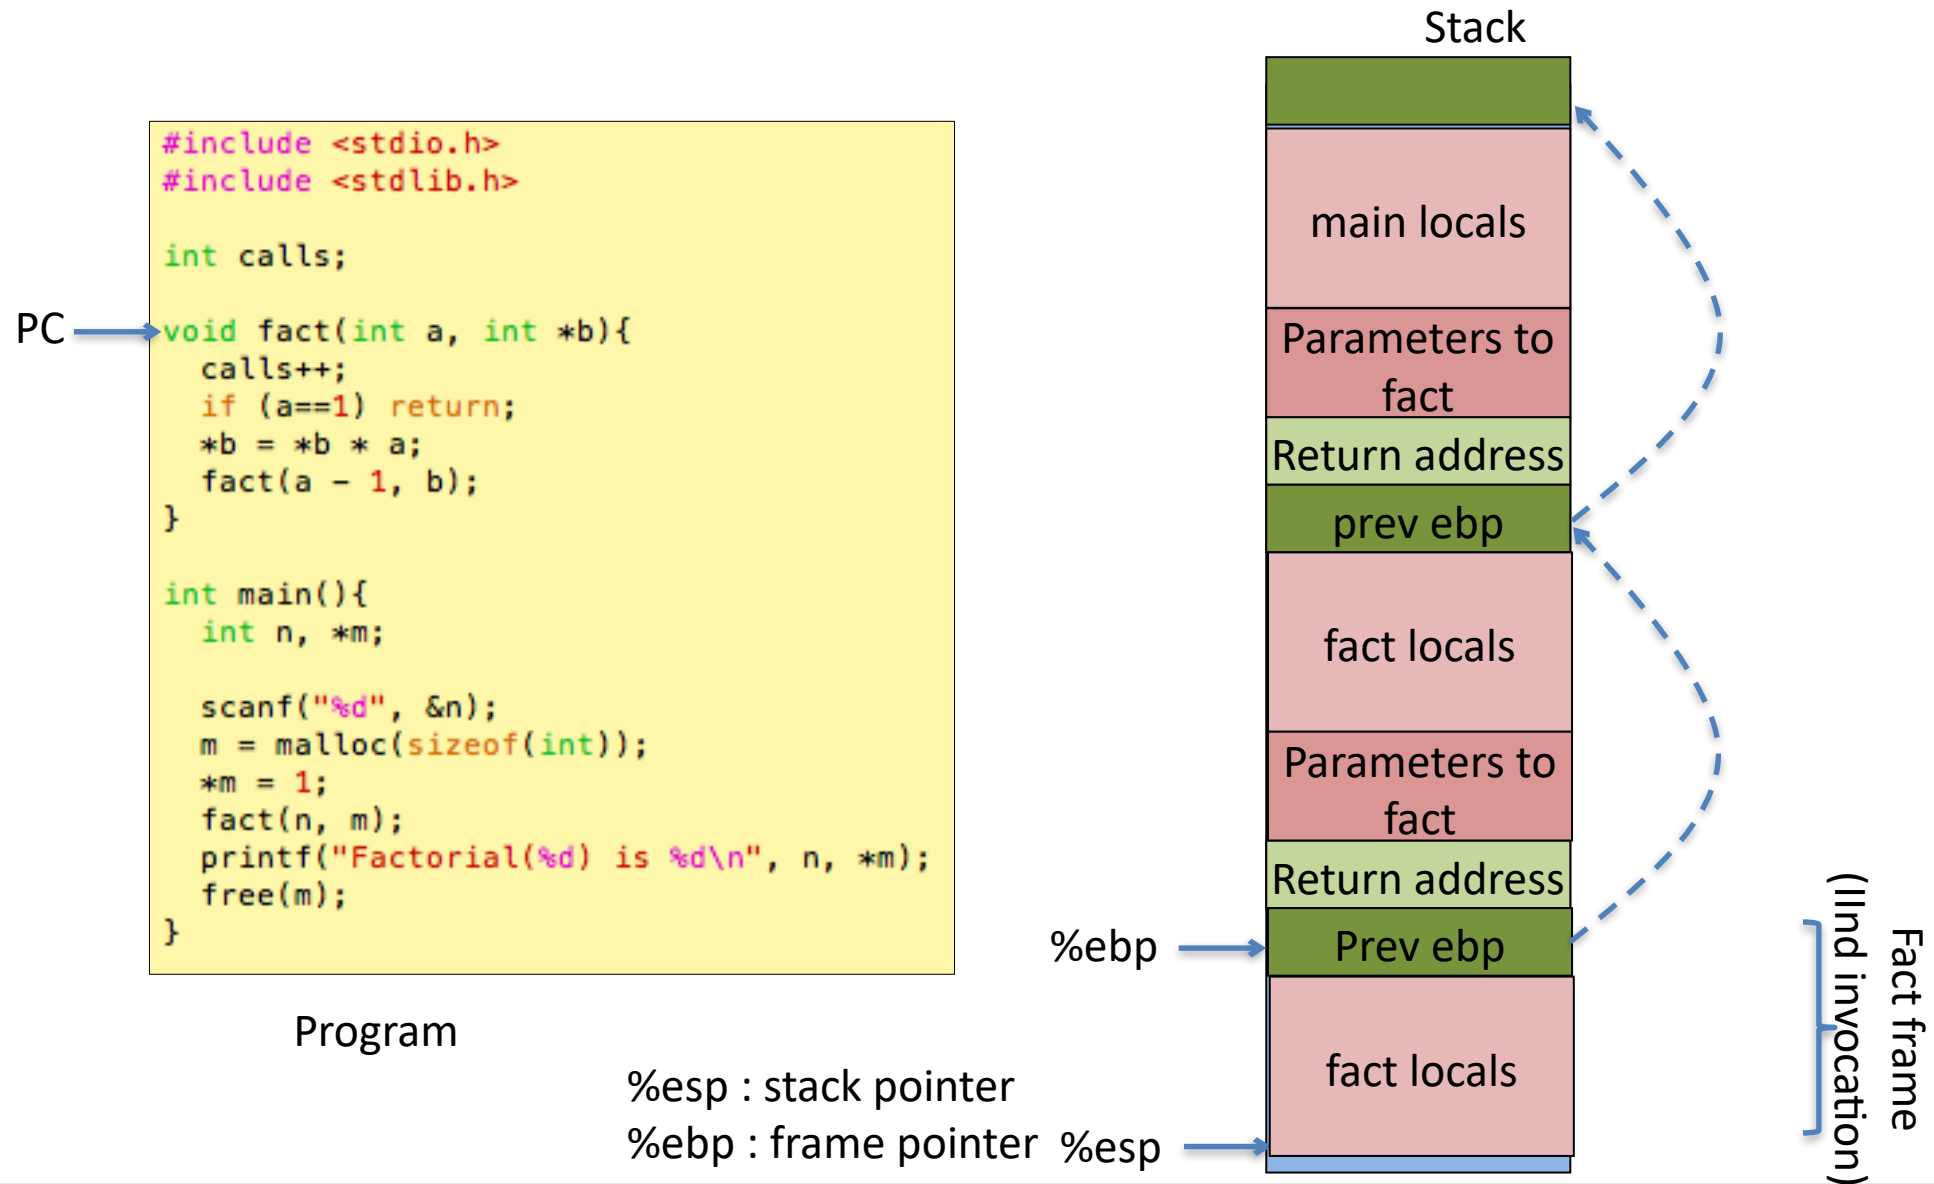
# Stack Frames



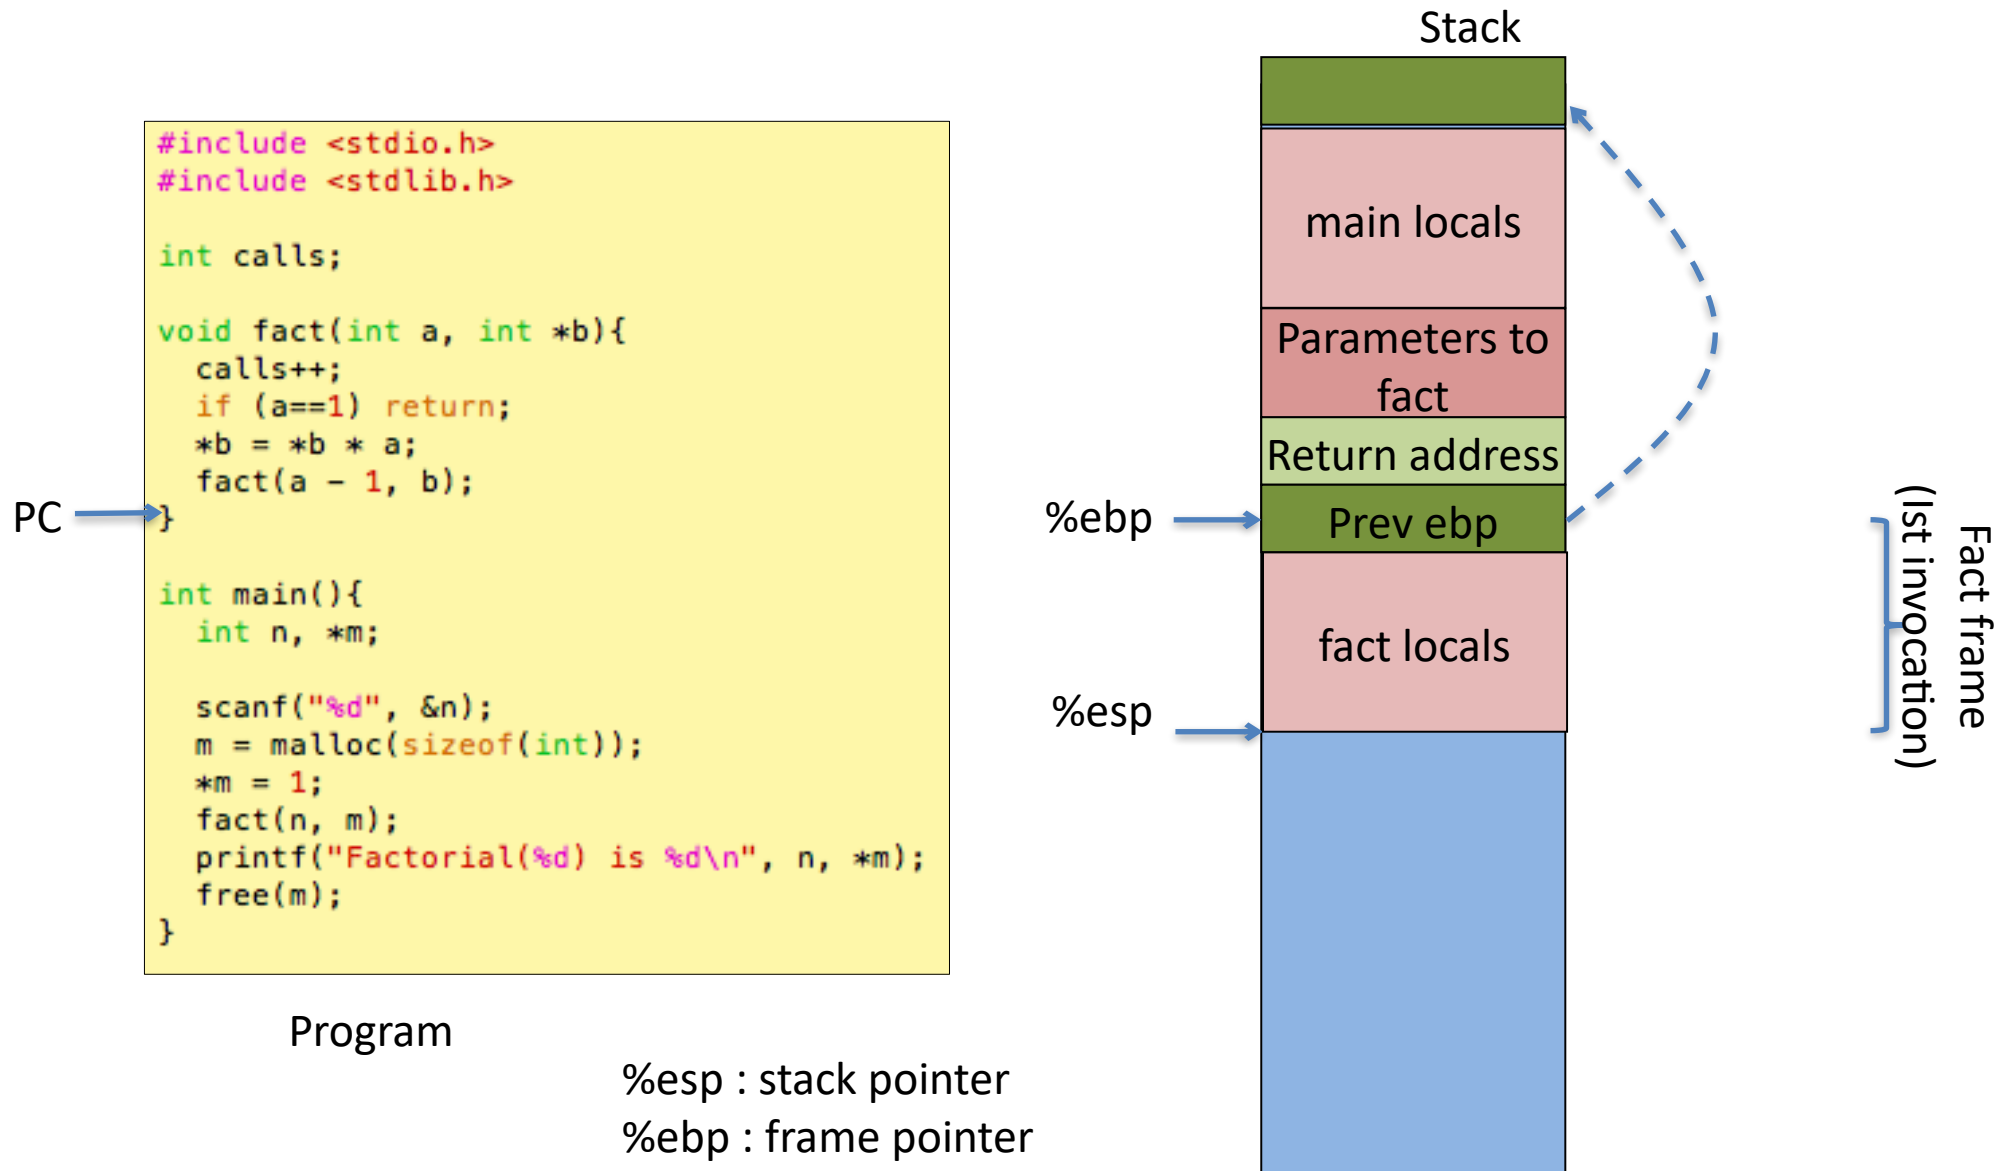
# Stack Frames



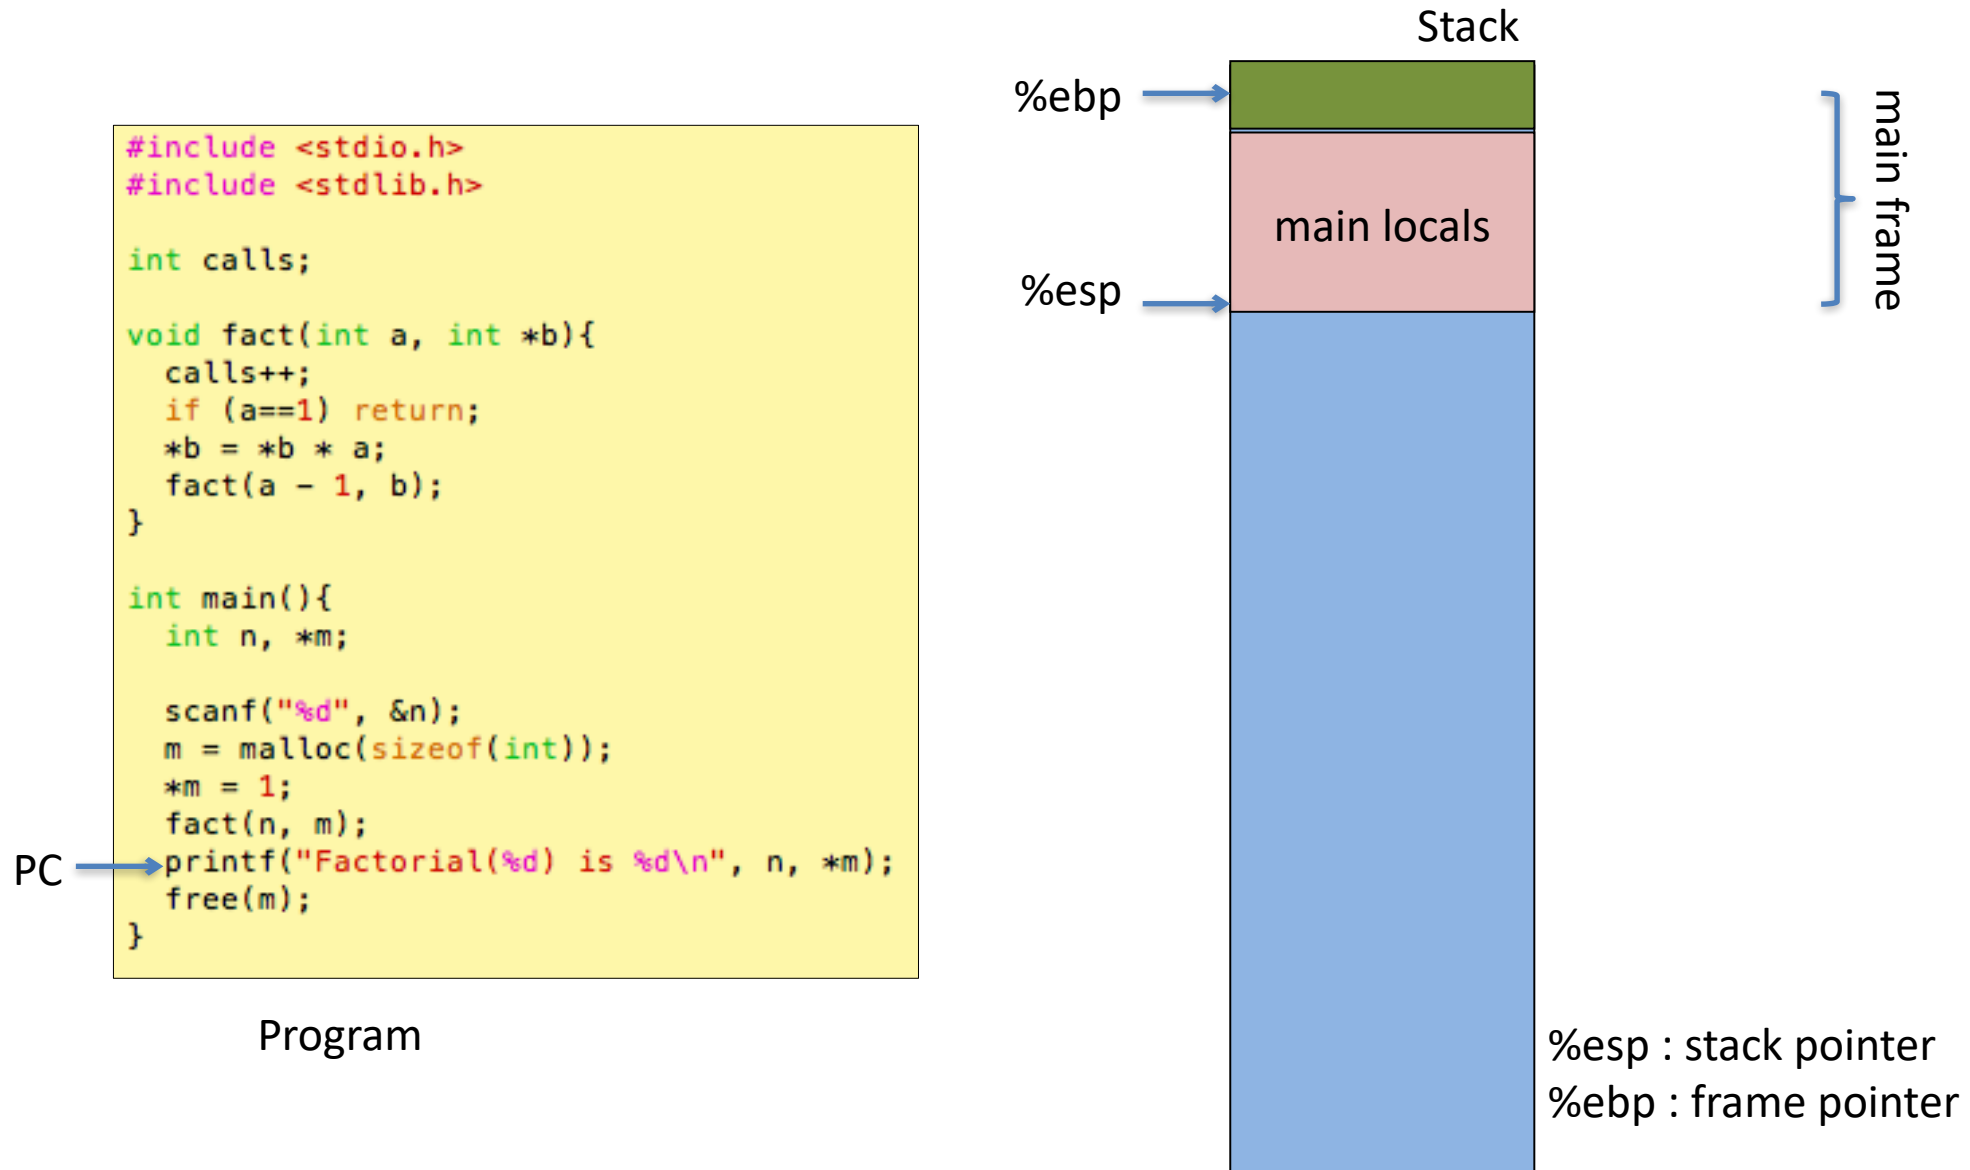
# Stack Frames



# Stack Frames



# Stack frames



# Points to Ponder

How and who passes command line arguments to the process?

